

Nest.js for Angular Developers

Enterprise Development with Nx Monorepo (Elvis)

@nxpatterns - <https://x.com/nxpatterns>
2025-01-29 | 19:09 | v0.5.7

Contents

| | | |
|----------|--|-----------|
| 1 | Nx Enterprise: Angular Expert's Guide to Nest.js | 5 |
| 1.1 | The Holy Trinity of Enterprise Development | 5 |
| 1.1.1 | Angular: The First Apostle of Type-Safe Frontend Salvation | 5 |
| 1.1.2 | NestJS: Modernizing Legacy Data Without Breaking Faith | 5 |
| 1.1.2.1 | How to Pitch NestJS to Your CTO? | 6 |
| 1.1.2.2 | Making Backend Developers Question Their Life Choices | 7 |
| 1.1.2.3 | Why not use Angular Server Rendering? | 7 |
| 1.1.2.4 | From Angular to Nest.js: Same But Different | 7 |
| 1.1.2.5 | Field and Battle Testing | 8 |
| 1.1.3 | Nx: The Holy Spirit of Enterprise Scalability | 8 |
| 1.1.3.1 | Why not use Angular CLI? | 8 |
| 1.1.3.2 | Why Nx? | 9 |
| 1.1.3.3 | Cross-Framework Library Re-Usability | 9 |
| 1.2 | Enterprise Schema Design: Angular & Nest.js | 10 |
| 1.2.1 | JSON Schema | 10 |
| 1.2.1.1 | JSON Schema Validation | 11 |
| 1.2.1.2 | Why JSON Schema? Think Enterprise | 11 |
| 1.2.2 | Schemas and Schematics | 11 |
| 1.2.3 | Differences in Schematics: For Angular Experts | 12 |
| 1.2.4 | Version Clarity: Angular & Nest.js Nomenclature | 13 |
| 1.3 | When Angular Met Nest.js: A Developer's Love Story | 14 |
| 1.3.1 | The First Date: Angular's Perspective | 14 |
| 1.3.2 | The Perfect Match: Nest.js Joins the Party | 15 |
| 1.3.3 | Nx Workspace: Angular and Nest.js | 17 |
| 1.4 | Install Recommended Software | 18 |
| 1.4.1 | VSCode | 18 |
| 1.4.1.1 | Download and Install | 18 |
| 1.4.1.2 | Register <code>code</code> as Shell Command | 19 |
| 1.4.1.3 | Why VSCode for Initial Development? | 19 |
| 1.4.2 | Node.js | 20 |
| 2 | Create a New Nx Workspace (Angular Setup) | 22 |

| | | |
|----------|--|-----------|
| 2.1 | Give Your Workspace a Name | 22 |
| 2.2 | Confirm Download & Workspace Installation Start | 23 |
| 2.3 | Choose Your Stack: Angular | 23 |
| 2.4 | Multiple Projects Setup: Integrated Monorepo | 24 |
| 2.5 | Set Your Initial Application's Name | 24 |
| 2.6 | Choose Your Bundler | 24 |
| 2.7 | Set Default Stylesheet Format | 25 |
| 2.8 | SSR and SSG/Prerendering Setup | 26 |
| 2.8.1 | Server-Side Rendering (SSR) | 26 |
| 2.8.2 | Static Site Generation (SSG/Prerendering) | 30 |
| 2.8.3 | Challenges | 30 |
| 2.8.4 | Architect's Notes | 31 |
| 2.8.5 | SSR & SSG in Nx: Decision Time | 33 |
| 2.9 | Choose Your Test Runner (E2E) | 33 |
| 2.9.1 | Key Principles | 33 |
| 2.9.2 | Challenges | 34 |
| 2.9.3 | Decision Time | 35 |
| 2.10 | CI Provider Section | 35 |
| 2.11 | Remote Caching | 36 |
| 2.12 | Final Confirmation | 37 |
| 3 | Exploring the Nx Angular Workspace | 38 |
| 3.1 | Open Your Workspace in VS Code | 38 |
| 3.1.1 | VSCode UI Interface Layout and Terminology | 38 |
| 3.1.2 | Checkpoint: 001-elvis-after-creation | 39 |
| 3.2 | Install Recommended VSCode Extensions | 39 |
| 3.2.1 | Create .vscode/extensions.json File | 39 |
| 3.2.1.1 | Windows: Creating Dot-Folders | 40 |
| 3.2.2 | Add Recommendations to extensions.json | 40 |
| 3.2.3 | Restart VS Code | 41 |
| 3.2.3.1 | Refreshing Window without Restart | 41 |
| 3.2.4 | Checkpoint: 002-elvis-after-recommended-extensions | 42 |
| 3.3 | Angular Setup Review | 42 |
| 3.3.1 | Serve the Angular Application | 42 |
| 3.3.2 | Checkpoint: 003-serve-the-angular-application | 47 |
| 3.3.3 | Setting Up a Clean Workspace | 48 |
| 3.3.4 | Checkpoint: 004-setting-up-a-clean-workspace | 50 |
| 4 | NestJS Setup | 51 |
| 4.1 | Preparing Nx for NestJS | 51 |
| 4.2 | Creating NestJS App Using Nx Generator UI | 51 |
| 4.2.1 | Application Type (@nx/nest) | 52 |
| 4.2.2 | Working Directory | 52 |
| 4.2.3 | Basic Configuration | 53 |
| 4.2.3.1 | Directory | 53 |
| 4.2.3.2 | Frontend Project | 53 |
| 4.2.3.3 | Linters | 53 |
| 4.2.3.4 | Name | 53 |

| | | |
|----------|---|-----------|
| 4.2.4 | Extended Configuration Settings | 54 |
| 4.2.4.1 | E2E Test Runner | 54 |
| 4.2.4.2 | Parser Options | 54 |
| 4.2.4.3 | Strict Mode | 55 |
| 4.2.4.4 | Tags | 55 |
| 4.2.4.5 | Unit Test Runner | 56 |
| 4.2.5 | Completing Configuration | 56 |
| 4.2.6 | Checkpoint: 005-create-a-nestjs-app | 56 |
| 4.3 | NestJS Setup Review | 57 |
| 4.3.1 | Serving the NestJS Application | 58 |
| 4.3.2 | Checkpoint: 006-serve-the-nestjs-application | 60 |
| 5 | Architectural Parallels | 61 |
| 5.1 | A Natural Partnership: Angular & NestJS | 61 |
| 5.2 | Decorators & Reflect Metadata API | 62 |
| 5.2.1 | The Decorator Decision: A Pivotal Moment in Angular’s History . . . | 63 |
| 5.2.2 | From Boilerplate to Elegant Architecture: Angular’s Decorator Journey | 63 |
| 5.2.3 | Angular’s Strategic Evolution | 64 |
| 5.3 | NestJS: The Rise of an Enterprise Node.js Framework | 65 |
| 5.3.1 | NestJS’ Strategic Evolution | 65 |
| 6 | Upgrade NestJS to v11+ in an Nx Workspace | 67 |
| 6.1 | Checkpoint: 007-upgrade-nestjs | 70 |
| 6.2 | The “ncu -u” Upgrade Experiment | 70 |
| 6.2.1 | Beyond Migration: Tracking Critical NestJS Updates | 74 |
| 6.2.2 | Checkpoint: 007A-upgrade-nestjs-experiment-with-ncu | 74 |
| 6.3 | From Nest CLI to Nx: Upgrade & Migration Patterns | 74 |
| 6.4 | Re-Check Versions & Patches | 76 |
| 6.5 | Production Readiness: Battle-Testing Our Upgrade | 77 |
| 6.5.1 | Current State Analysis | 77 |
| 6.5.2 | Official vs. Custom Migrations | 78 |
| 6.5.3 | Our Enterprise Migration Manifest | 78 |
| 6.5.4 | Necessary Steps: The Security-First Approach | 78 |
| 6.5.4.1 | Remove Old Webpack Configuration | 78 |
| 6.5.4.2 | Integrate nest-cli.json | 79 |
| 6.5.4.3 | Update NestJS’ project.json | 79 |
| 6.5.4.4 | The Surgical Strike: Security Audit | 80 |
| 6.6 | Checkpoint: 008-nestjs-basics | 82 |
| 7 | NestJS: When Angular Developers Go Full-Stack | 83 |
| 7.1 | Basic NestJS Concepts for Angular Developers | 83 |
| 7.1.1 | The NestJS Module: A Familiar Face | 83 |
| 7.1.2 | Controllers | 84 |
| 7.1.3 | Providers | 86 |
| 7.1.4 | Middleware | 87 |
| 7.1.4.1 | Middleware Outside the Box | 89 |



© x.com/nxpatterns - 2025

This work is licensed under the **Creative Commons Attribution-NonCommercial 4.0 International License**. To view a copy of this license, visit: creativecommons.org/licenses/by-nc/4.0/

Additional Terms:

- You are free to share and adapt this work by giving appropriate credit but you may NOT sell, resell, or commercially distribute this document itself.
- To use this document in commercial training materials or paid content, just tag ([@nxpatterns](https://x.com/nxpatterns)) on X-Platform and state your intended use. Permission will be granted immediately.

Disclaimer and Professional Notice

The concepts, architectures, and implementations described in this documentation are intended for use under the guidance of experienced software architects and professionals. While the presented approaches and patterns have been thoroughly tested in production environments, their successful implementation requires:

- Oversight by experienced software architects
- Proper evaluation of specific use cases and requirements
- Thorough understanding of architectural implications
- Appropriate risk assessment and mitigation strategies
- Consideration of organizational context and constraints

THIS DOCUMENTATION AND THE ACCOMPANYING CODE ARE PROVIDED “AS IS” WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESS OR IMPLIED, INCLUDING, BUT NOT LIMITED TO:

- The documentation and code examples are provided for educational and informational purposes only
- No warranty or guarantee is made regarding the accuracy, reliability, or completeness of the content
- The author(s) assume no responsibility for errors or omissions in the content
- The author(s) are not liable for any damages arising from the use of this documentation or implementations derived from it. Implementation decisions and their consequences remain the sole responsibility of the implementing organization
- Success in one context does not guarantee success in another; proper evaluation is required

To minimize risks and ensure successful implementation:

- Conduct thorough architecture reviews before implementation
- Engage experienced software architects throughout the project lifecycle
- Implement comprehensive testing strategies
- Maintain proper documentation of architectural decisions
- Consider scalability, maintenance, and long-term implications
- Establish proper governance and review processes
- Ensure adequate team training and knowledge transfer

Implementation of the described architectures and patterns requires substantial expertise in:

- Enterprise software architecture
- Monorepo management and tooling
- Modern development practices and toolchains
- System design and integration
- Performance optimization and scaling
- Security best practices
- DevOps and CI/CD processes
- Azure, AWS, Google Cloud, Hetzner Cloud, or other cloud platforms

All trademarks, logos, brand names, and product names mentioned in this document are the property of their respective owners. This includes but is not limited to Angular, Nx, Nest.js, TypeScript, and Node.js. This document is for educational purposes only. Any references to protected intellectual property are made under fair use for the purpose of technical documentation and training. The author makes no claim to ownership of any third-party intellectual property referenced herein. All rights remain with their respective owners.

1 Nx Enterprise: Angular Expert's Guide to Nest.js

1.1 The Holy Trinity of Enterprise Development

Having worked in enterprise software architecture, **you'll often encounter data structures that reflect their historical evolution** - shaped by numerous stakeholders, each contributing their unique terminology.

Let's just say it's an interesting archaeological expedition

Now, imagine the scenario: A new app needs to be developed, integrating with the existing API infrastructure. The deadline? Yesterday, of course! It needs to showcase cutting-edge tools, programming languages, and modern interfaces. Why? Perhaps a new CEO or government official is eager to demonstrate innovation in their next public appearance. After all, staying ahead of the competition is the name of the game!

1.1.1 Angular: The First Apostle of Type-Safe Frontend Salvation

Speaking of frontends; **What could possibly outshine Angular?** Especially when operating in German-speaking Europe? **Surprisingly, Angular is the framework that gets executive approval fastest** as executives recognise it as a **“modern interface solution”** - it's familiar to them, they've heard of it repeatedly, and chances are there are already several Angular applications running somewhere in the organisation.

Would you like to enhance your Angular knowledge? I've prepared a chapter at the end of the document for you with valuable resources.

1.1.2 NestJS: Modernizing Legacy Data Without Breaking Faith

Now that we've settled the frontend question, what about the dreadful data that has evolved organically over years?

If Murphy's Law applies (and it always does), you might find database tables and columns in Low German dialect, leaving international team members completely lost - or as we say in German, they “understand only train station” (*verstehen nur Bahnhof* - an idiom for understanding absolutely nothing). Even as a native German speaker, I sometimes feel lost too - Low German can be as cryptic as hieroglyphics.

Understanding individual elements isn't always sufficient, as context can dramatically alter their meaning. Mastering domain-specific terminology (including considerations like semantic color schemes) is essential for successful modernization.

Now that we've “diplomatically acknowledged” that the existing data structure needs improvement, let's focus on modernizing this legacy system.

1.1.2.1 How to Pitch NestJS to Your CTO?

Best case scenario: You have a visionary CTO. Worst case scenario: You have an accountant in CTO's clothing¹.

- If we position it as a backend, the company's backend team (our API providers) might either feel threatened or suddenly develop "concerns" that get us ejected from the project.
- If we try selling it as a "frontend backend," the CTO's disciples will immediately interrupt with "*Why not use Angular's server rendering features?*" (which, admittedly, becomes a valid question from Angular 19+ onward).
- And if we dare to mention "middleware" at an unfavourable moment, the CTO will likely ask, "*What's that?*" - and we'll be forced to explain it as "a layer between the frontend and backend that handles business logic, data transformation, and validation." If he asks "Do we really need that?" - the answer is "Yes, my Lord," especially if the company size is at least 1999 employees or larger.

I usually call it "**middleware**" - fully aware that NestJS has its own different interpretation and implementation of this term². *My sincere apologies to the NestJS team, but perhaps they might consider renaming their middleware concept... for my sake?*

From an architectural perspective, I prefer using NestJS to implement "**Separation of Concerns**". This pattern helps maintain clean code boundaries by keeping business logic and data transformations in the middleware layer rather than the frontend.

Security best practices recommend implementing validation on both frontend and backend layers. While frontend validation improves user experience by **providing immediate feedback**, it can be circumvented using browser developer tools. Therefore, robust middleware validation through NestJS serves as a critical security measure (I'll demonstrate these security implications in detail later).

When **deploying Angular and NestJS**, I prefer using two different domains. If that's not feasible (*your accountant-turned-CTO might consider €12 per year an "unnecessary infrastructure expense"*), I'll use subdomains instead.

Under no circumstances would I advocate for hosting everything on the same domain with different ports (unless, of course, I'm developing locally on my own machine). This distinction is crucial, and I'd be delighted to elaborate on its importance at a later time.

This middleware is only accessible by authorized frontends that possess valid tokens (more on that later). *I've even offered some CTOs to personally sponsor these €12 per year - though that suggestion wasn't particularly well-received.*

¹ This chapter is dedicated to [@kammysliwiec](#), the creator of NestJS. His dedication to the project and the community is truly admirable.

² Middleware in NestJS is a function which is called before the route handler. Middleware functions have access to the request and response objects, and the `next()` function. NestJS is based on Express and these concepts were adopted directly. Learn more here: expressjs.com/en/guide/using-middleware.html. While serving different environmental contexts, NestJS middleware can be conceptually compared to Angular's HTTP interceptors, as both handle cross-cutting concerns in their respective request pipelines. However, this comparison might raise some eyebrows since NestJS also offers its own interceptors docs.nestjs.com/interceptors.

1.1.2.2 Making Backend Developers Question Their Life Choices

Nest.js stands as a comprehensive backend framework capable of replacing most conventional backend solutions. While its TypeScript foundations might raise eyebrows among traditional backend developers³, its capabilities are undeniable.

This is precisely why we present **Nest.js** to “**traditional**” **backend developers** as an intermediate layer between our frontend and the “**core**” **backend**. This approach excels in scenarios requiring data modeling, internationalization, and modernized business logic handling.

While Nest.js often proves indispensable in projects dealing with historically (and sometimes hysterically) evolved data structures, positioning it as a middleware layer not only makes tactical sense but frequently becomes **a technical necessity**.

1.1.2.3 Why not use Angular Server Rendering?

The legitimate question “Why not use Angular Server-Side Rendering (SSR)?” deserves an answer: As Software Architects, we need to extensively field-test and validate “Angular SSR”.

While Angular v19+ brings some stability (please don’t stone me, but I personally had sub-optimal experiences with earlier versions), we already have NestJS as middleware (in my terminology) thoroughly field and battle-tested.

It’s simply not acceptable for an architect to build a modern, daring stadium that collapses during a live game. Nobody wants to be quoted in the tabloid headlines (especially not in “Bild”). That would be career suicide - you’d be relegated from designing architectures to carrying sandbags.

1.1.2.4 From Angular to Nest.js: Same But Different

When transitioning **from Angular to Nest.js**, you’ll notice several similarities. **Both** frameworks are built on **TypeScript**, and both use **decorators** to define classes and methods.

The module system in Nest.js will feel familiar to Angular developers, though with some key differences. Controllers in Nest.js might remind you of the old Angular.js days, and while they serve a similar routing and request handling purpose as Angular components, they’re fundamentally different. Think of them as the entry points for your HTTP requests rather than UI elements.

Currently, Nest.js doesn’t have an equivalent to Angular’s standalone components - though who knows what Kamil (Nest.js creator) has planned for the future.

While these architectural parallels appear complex at first, we’ll explore all these details hands-on as we create a Nest.js application within our Nx monorepo. After all, learning by doing is the best approach.

³ The debate around “proper” backend languages often sees JavaScript/TypeScript facing skepticism from developers who favor languages like Go, Rust, Java, or C. Interestingly, even Python, despite its AI dominance, hasn’t fully established itself in traditional backend development. While frameworks like Django (currently the de facto standard), Flask, CherryPy, and TurboGears have made notable attempts, Python’s backend adoption remains primarily in specialized domains like data science and machine learning. Its widespread backend adoption would likely require significant hardware optimizations at the processor level.

1.1.2.5 Field and Battle Testing

How does one properly field and battle-test?

Simple: Rent a cloud server from Hetzner⁴, create an app with Angular Server-Side Rendering (without a firewall), then post on Reddit or the dark web claiming, “*I’ve built the most secure frontend web app - who dares to challenge it?*”

If it survives six months without firewalls (but with OS updates), it might be production-ready.

But seriously, **security testing requires academically recognized metrics**. I haven’t had the chance to thoroughly evaluate Angular v19+ yet - I’m still learning myself. The Angular team has done a masterful job here, and I tip my hat to their achievement.

And **testing**, whilst not quite a science, is certainly an art form: Unit testing, Integration testing, End-to-End testing (E2E), Security and Penetration testing, Usability and Accessibility testing, Responsiveness testing, Local testing, Pipeline testing, Background offline testing (during development as pre-commit and pre-push rules), Test-Driven Development (TDD), and Behaviour-Driven Development (BDD - which is essentially TDD done right).

I realise I might be turning this quick start guide into a doctoral thesis. But I say, when it comes to testing, it’s rather easy to fall down the rabbit hole.

And thus, we have discussed two Apostles of our Holy Trinity. But what of the third one? What divine powers does Nx bring to our enterprise salvation?

1.1.3 Nx: The Holy Spirit of Enterprise Scalability

I’ve been using Nx since its inception⁵. When we want to create enterprise workspaces that offer a proven schema from day one (and through my own enhancements, provide infinite scalability and true re-usability - more on that later), Nx is ideal.

You know naming is hard, especially with multiple participants pulling in different directions. Well, surprisingly, I have a secret formula here: **ELVIS**. When I propose this name, no one rejects it. The business departments and domain experts are extremely creative, quickly transforming ELVIS into “electronic files”, “super builder”, “enterprise live information system”, and so on.

I’ve been building enterprise workspaces named ELVIS since around 2008, and since 2018, I’ve been doing it with great joy using Nx. So, if you’re in a company with a workspace/framework/solution/product/... named ELVIS, there’s a high probability that was my doing.

1.1.3.1 Why not use Angular CLI?

The reason is simple: the strong, omnipresent React community. But jokes aside, when a smarty-pants department head jumps up during my ELVIS presentation declaring “I don’t

⁴ **Hetzner.com** is a German cloud provider known for its affordable prices and reliable services. It’s a popular choice among European developers and it is also my preferred choice - they’re an environmentally conscious company where I train my models cost-effectively. **Full disclosure:** I have no business relationship with Hetzner and am not affiliated with them.

⁵ I just checked - I actually started using Nx a year after its creation, which proves my cautious approach to adopting new technologies. See for yourself: <https://api.github.com/repos/nrwl/nx>

like Angular”, the response is immediate: “I love React too (*well, actually I just like her - it’s not true love*) which is why you’re welcome to develop your apps within ELVIS using React.”

Usually, they sit down with a smile, but then another one springs up: “Can we use the libraries you’ve developed for Angular in React?” And if you know how, the answer is an immediate yes.

One of the main reasons we don’t use Angular CLI is that large enterprises (typically operating internationally) want to bundle their applications and products for **specific markets** and offer them as a **unified software suite**. There are countless **existing apps**, and **they’re not all necessarily Angular applications** - there’s **everything** in the mix.

When **establishing our Nx Mono Repo approach** (*which we’ve named ELVIS*), existing applications (Angular, Vue, React, whatever framework you prefer) need to **become ELVIS-compatible**.

We typically prepare **automated import mechanisms** to manage this process. For these mechanisms to work effectively, TypeScript is a technical prerequisite. (*The comprehensive benefits and reasoning behind this choice will be explored in a dedicated chapter.*)

This necessity often leads to the **delicate task** of requesting all participating departments, products, and sub-products **to migrate their JavaScript applications to TypeScript**. While this typically triggers **an avalanche of reactions** across the organization, it’s a **necessary transition** for achieving our architectural goals.

If the applications are already written in TypeScript, the integration becomes significantly simpler. And if teams have developed their apps in a current or previous Nx Mono Repo setup, the process becomes even more straightforward.

1.1.3.2 Why Nx?

While I won’t market Nx - and frankly, it doesn’t need marketing - its practical value for architects is clear: The workspace schema naturally supports fractal structures (as Mandelbrot would appreciate), making the internal logic instantly recognizable to all stakeholders.

The library management enables genuine re-usability (more on real versus fake re-usability later), while the CLI and plugin architecture allow for extensive automation and customization.

From an architect’s perspective, Nx brings three critical capabilities:

- Enforced structural consistency across large-scale projects
- Granular control over module boundaries and dependencies
- Scalable build and test orchestration

For detailed technical reference and implementation guidance, consult the comprehensive Nx documentation: <https://nx.dev/>

1.1.3.3 Cross-Framework Library Re-Usability

When introducing Nx with Angular, we occasionally encounter raised eyebrows. React developers - as previously mentioned - often inquire about Angular libraries’ compatibility with React. The answer is an unequivocal “yes.”

Beyond this documentation repository, I plan to establish an ELVIS repository that will evolve step-by-step into an enterprise-ready Nx monorepo workspace. Detailed examples will follow as we progress through:

- Starting with a bare Nx framework
- Integrating Angular and Nest.js
- Developing our reusable libraries (is in another documentation)

For those eager to dive in, let's clarify the path to **cross-framework harmony**. You can achieve cross-framework library compatibility when any of these conditions are met.

- Core logic in TypeScript without framework dependencies
- Thin framework-specific wrappers for Angular, React, or others
- Shared interfaces for consistent data contracts
- Framework-agnostic state management
- Common utility functions and services

And I've developed a sixth approach: a special kind of Proxy Pattern. Not the traditional OOP (Object-Oriented Programming) proxy, but rather a translator between apps that makes the whole interaction transparent and straightforward (more on this later).

1.2 Enterprise Schema Design: Angular & Nest.js

Before diving into implementation, let's clarify some fundamental concepts that often cause confusion. Understanding these will be essential for the chapters ahead.

1.2.1 JSON Schema

JSON Schema provides a declarative vocabulary to validate JSON documents. Think of it as a strongly-typed contract that defines the expected structure, types, and constraints of your data.

Example: For a person object, a JSON Schema might specify:

- **name:** Required string, e.g. "Noah of Ancient Mesopotamia"
- **age:** Required number between 0-1000 (Don't be mad at me, think of Noah)
- **address:** Optional nested object with its own schema

In **Nx** monorepo contexts, JSON Schema validates crucial configuration files like `nx.json`. It ensures that workspace settings (`namedInputs`, `targetDefaults`, `plugins`, `generators`, etc.) maintain structural integrity and type safety.

Since `nx.json` now serves as the central configuration hub, the schema acts as a guardrail, preventing misconfigurations in project definitions, task executions, and dependency management.

1.2.1.1 JSON Schema Validation

The following diagram illustrates the principle of JSON Schema Validation (For more details, please refer to the official documentation at: json-schema.org):

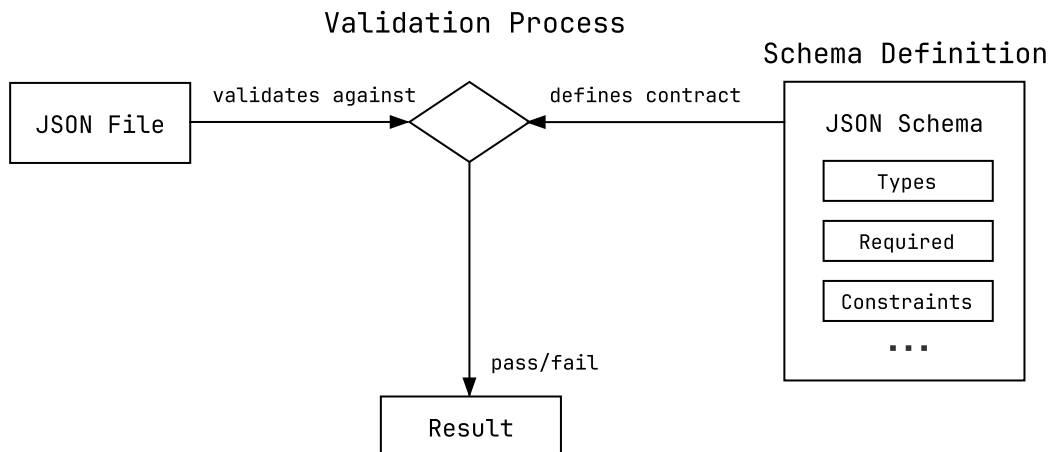


Figure 1: JSON Validation Process Overview

1.2.1.2 Why JSON Schema? Think Enterprise

Remember that awkward moment when you tried to explain to your boss *why all your apps look different*? Well, that's exactly what we're preventing here!

Nx schema isn't just another buzzword - it's your secret weapon for building enterprise-grade applications that play nicely together. Think of it as a strict but fair bouncer at an exclusive enterprise club:

- Every app built following the schema gets a VIP pass (**auto-recognition by Nx**)
- No more “but it works on my machine” drama (**consistent configurations**)
- Future-proof your apps for the ultimate enterprise party (**suite integration**)

Both Nx and Angular use JSON files for project/app configurations that are based on JSON Schema. (We'll examine these schemas in detail in another document.)

1.2.2 Schemas and Schematics

Following our deep dive into **JSON Schema**, let's connect the dots with **Schematics**: Schematics are essentially tools, **powerful tools**. They automate and standardize development workflows.

💡 Schematics in a Nutshell

Imagine you're running a high-end restaurant. **JSON Schema** acts as your health inspector ensuring ingredients meet standards, while **Schematics** is your master robot chef who executes recipes flawlessly.


In **Nx terms**, running a generate command for a new library is like ordering your robot chef to prepare a new signature dish - it creates all necessary ingredients (**files**), updates the menu (**configuration**), and even sets up the kitchen layout (**project structure**) automatically.

One simple command, and your workspace transforms like a well-orchestrated kitchen during service time!

These clever transformations are powered by Schematics under the hood, which follows your JSON Schemas like a precise recipe book - ensuring every ingredient (file), measurement (configuration), and cooking step (code generation) meets your exacting standards, like a personal sous chef who never gets tired, never makes mistakes, and always delivers consistent results.

1.2.3 Differences in Schematics: For Angular Experts

Let's examine these two (simplified) structures:



On the left, we have an Angular app generated using Angular/CLI in a directory named “my-angular-app”. In our case, the directory name is the app name. However, if we were to create

libraries for this app outside of `src` folder, this directory would quickly evolve into a small “workspace” in architectural terminology.

On the right, we have a workspace (established as such from the beginning) that we named `elvis-nx-workspace` (when I create an Nx workspace later, I’ll simply name it `elvis`), which contains multiple Angular apps. Both apps have their own `project.json`, and there’s no `package.json` in their directories since the workspace itself **has a single one**. If we were to create libraries, any Angular app could use these libraries. (The orchestration here is slightly different; we’ll learn about the principle of *infinitely scalable fractals* and *true reusability* of global libraries later - libraries that can be developed *without interfering with each other*. This will probably be covered in another document or book, we’ll see.)

Nx has elegantly streamlined Angular’s sophisticated blueprint, optimizing it for multi-project architectures while preserving its core strengths. We’ll explore this architectural evolution - particularly the transition from `angular.json` to `project.json` - when we dive into advanced workspace patterns and set up our Nx workspace together.

1.2.4 Version Clarity: Angular & Nest.js Nomenclature

When searching for “Nest.js”, you’ll find “NestJS” (for SEO purposes). The official website shows it as “Nest.js” and uses nestjs.com domain. In this document, we’ll use **Nest.js**, following the official homepage nomenclature.

Using Synonyms

While we can use **Nest.js** and **Nest** (or **NestJS**) synonymously, we can’t do the same with **Angular** and **Angular.js**^a.

Because with AngularJS, we’re referring to the first initial version of the framework (version 1.x), which was released in ~2010. Angular (referred to as Angular 2+) is a complete rewrite of AngularJS, released in ~2016. **Angular.js and Angular are fundamentally different frameworks**, so it’s crucial to distinguish between them.

^a Angular experts are already familiar with the difference between Angular.js and Angular.

When we say Angular today, we’re always referring to the modern framework (Angular 2+), not the older Angular.js. As of this writing, the latest Angular version is 19.0.5⁶.

⁶ See also: nx.dev/nx-api/angular/documents/angular-nx-version-matrix

1.3 When Angular Met Nest.js: A Developer's Love Story

1.3.1 The First Date: Angular's Perspective

When our Angular experts establish the architecture, they know we need a configuration file that defines our `https://what-ever-it-is/api` (depending on the environment).

After all, we typically want to access an API and fetch our data from there, even though **Angular itself doesn't care where the data resides**. Since this document was written for Angular experts, intended to provide them with a quick introduction to the Nest.js world (within Nx), this context is particularly relevant.

This “*what-ever-it-is*” is typically `localhost` for local development, but only if our backend is also available locally. If the backend resides elsewhere, the targets should be configured accordingly. This mechanism is called “**Proxy Environment Configuration**” - at least that's what I call it in my documentations, and I believe it hits the nail right on the head. The allowed targets (targets we're permitted to access) are defined within this file.

You've guessed it correctly - this file should follow community best practices for naming (typically called `proxy.conf.json`, though it could technically be named anything else, but shouldn't be), should be located in a specific place (best practice: in the root directory of the respective Angular app), and must follow a specific schema.

It's crucial **not to confuse** this “**Proxy Environment Configuration**” with the `localhost` that Angular starts by default when running `ng serve` or `nx run <my-app>:serve`. These are two distinct configurations, though both use `localhost`. For the frontend, we use `localhost` because I've started my Angular app locally, and I've set `localhost` as targets in my Proxy Environment Config because, coincidentally, my backend is also available locally.

In our case, we want to set up an Enterprise Nx Mono Workspace (which we've named ELVIS), and this is where the Nx perspective comes into play.

Let's summarize what we've learned about how things look for Angular experts who have developed standalone Angular apps without Nx, versus how the world looks in an Nx Workspace. But before the brief summary, an important note: For the Angular app, it's not enough to just have the correct targets in `proxy.conf.json` - the file must also be referenced⁷ properly. E.g. for local dev-servers:

For standalone Angular apps:

- Uses Angular Dev Server Proxy Configuration Schema
- Schema defined in [@angular-devkit/build-angular](#) package
- Configuration in `angular.json` references the proxy config

⁷ We must know these schemas precisely because even though Nx typically adopts Angular schemas one-to-one, it makes modifications or omits certain elements. (I'll revisit this in another document/book; currently, I have the example of `vendorChunk`, which was marked for removal in Angular source code, but Nx removed it even before Angular did). This means if you're reading an older Angular book and want to simplify/optimize debugging on a Development Server, Nx might follow different guidelines. This type of information rarely appears in books - you need to read the developer comments on GitHub yourself, in this case under **build-system-migration.md**.

For Angular apps in Nx mono repo:

- Uses Nx Dev Server Proxy Configuration Schema
- Schema defined in [@nx/angular](#) package
- Configuration in project.json references the proxy config
- Same schema structure (almost) but different implementation under the hood

Key insight: The configuration approach remains conceptually similar, but the implementation details and file locations differ between standalone Angular and Nx environments. Fortunately, we don't need to create many of these configuration files ourselves - they're generated by the Nx CLI while creating our apps. No worries, I'll also guide you through the few manual configurations required, explaining each step in detail.

1.3.2 The Perfect Match: Nest.js Joins the Party

Angular experts know this all too well: Angular has done everything right, but somehow it's still not working, and it's not even our fault. Yes, you guessed it - CORS errors it's like Angular and Nest.js are at a fancy developer party, but they can't dance together because the bouncer (the browser) is extremely picky about cross-origin protocols.

Imagine Angular, dressed in its best HTTP requests, trying to approach Nest.js:




Figure 2: CORS Party Sequence Diagram

This is the infamous Cross-Origin Resource Sharing (CORS) error - where the browser plays an overprotective guardian, ensuring domains can't freely share data unless explicitly allowed. It's a crucial security feature, preventing malicious websites from making unauthorized requests to your API.

“Don't worry, my dear Angular,” says Nest.js, “I'll add the proper CORS headers to my responses. Just tell me where you're from, and I'll put you on my VIP list. After all, security is important, but so is true love!”

And here comes the voice-over you know from the movies - the Architect. He must ensure that Nest.js gets to know its Origin. Alas, nobody knows him - this unsung hero who tirelessly orchestrates things in the background, firing the right arrows to ensure this love story reaches its happy ending.

Hold On: Setup Details Follow

Don't worry about the earlier mention of manual configuration - we'll cover Angular and Nest.js app setup in detail when we create them.

And so our diligent Architect performs the time-honored ritual of matchmaking. First, he creates a special list in a file called `localDevOrigins.ts` under Nest.js's `src` folder - think of it as Nest.js's dating preferences:

```
1 export const LocalDevelopmentOrigins = [  
2   'https://localhost:4200', // Formal address with a secure connection  
3   'http://localhost:4200',  // And the casual one  
4 ];
```

Then, in Nest.js's `main.ts` file, just before the grand opening (`await app.listen(port);`), the Architect whispers the sweet configuration of acceptance⁸:

```
1 import { LocalDevelopmentOrigins } from './localDevOrigins';  
2 // ...  
3  
4 app.enableCors({  
5   origin: [  
6     // other potential suitors  
7     ...LocalDevelopmentOrigins,  
8   ],  
9 });
```

⁸ Naming is always crucial - the file should be descriptive and follow a clear naming convention. Sometimes we also call it “local-dev-origins.ts” or “local-development-origins.ts” - this naming decision is made by the team in a shared ritual known as “Coding Style Guidelines”.

1.3.3 Nx Workspace: Angular and Nest.js

Now Angular and Nest.js stand face to face for the first time. Let's get their relationship status straight. To Nest.js, Angular is the great love, the irreplaceable wife, the so-called Origin. For Angular, however, Nest.js is simply the Proxy.





Figure 3: Relation

🐾 Why Proxy?

Well, it's kind of like a stand-in until the next one comes along, or maybe because it can be swapped out quickly and easily? Or perhaps because each environment has a different one? I'd better leave these philosophical questions to the philosophers, but I've a theory.

Both Angular and, naturally, Nx use the term `proxyConfig` in their schema. If you want to know exactly where, you'll need to look in the file: `@angular/cli/lib/config/schema.json`. In Nx, the schema is defined in `@nx/angular/src/builders/dev-server/schema.json` package. Please note that files, structures, and schemas may change in newer versions. This is just a snapshot in time.



🐾 Chronicles of Schema-Fu

As Master Oogway whispered: *Yesterday is history, tomorrow is a mystery, but your git history 🙄🙄. The schema you're looking for might not be the schema you need.*

Master Shifu answered: *True validation comes not from regex patterns, but from accepting that users will always find a way to break your schema. Yet you must still **find ways to protect the workspace**. A schema that bends will survive, but one that validates will prevail.*

⚠ Production Environments are Different

Keep in mind that serving an app locally can be entirely different from running it in a production environment (such as in the cloud where a web server might be running, or when your app is embedded in a content management system like Adobe Experience Manager)^a.

^a In my twenty years of software architecture experience, I thought I'd seen every possible variation imaginable - at least that's what I believed two years ago. Now I find myself in agreement with Socrates, albeit in a slightly different context: *The more I learn, the more I realise there's no limit to creative implementation approaches.*

1.4 Install Recommended Software

1.4.1 VSCode

1.4.1.1 Download and Install

We will use VSCode as IDE (Integrated Development Environment). If you don't have it installed, Download and install VSCode from <https://code.visualstudio.com/>.

Here's how the VSCode interface roughly looks like (simplified), Please remember names of these individual areas (that's the VSCode Interface terminology), as we'll refer to them later:




Figure 4: VCode Interface Overview

1.4.1.2 Register code as Shell Command

To seamlessly work with our development environment, we need to be **able to launch VS-Code from the terminal**⁹. This enables powerful workflows like:

- Opening project directories directly: `code my-project`
- Opening files at specific lines: `code file.ts:42`
- Managing VSCode from CI/CD pipelines
- Using VSCode as default Git editor etc.

Open the Command Palette

F1 Key Might Help too

If not mapped differently, the F1 key works across Windows, macOS and Linux. If not, then use the following shortcuts:

- Windows/Linux: `Ctrl+Shift+P`
- macOS: `(Cmd+Shift+P)`

Then type: **Shell Command:**

You should see in the Title Bar of the VSCode window a dropdown menu; select **Shell Command: Install 'code' command in PATH** and press Enter.




Figure 5: Registering `code` as Shell Command

Restart your terminal & verify installation:

```
1 code --version ↵
2 1.96.2 # or higher
```

1.4.1.3 Why VSCode for Initial Development?

As an architect, I fully support diverse development environments. However, when establishing a new enterprise-wide development initiative, it's strategic to start with a standardized, zero-cost solution that minimizes initial friction. **VSCode is an excellent starting point:**

- Zero licensing costs
- Wide enterprise acceptance (particularly in Microsoft-dominated environments)

⁹ The `code` command is usually auto-registered during VSCode installation on Windows. However, certain corporate policies may override this registration (don't ask me how & why).

- Robust extension ecosystem
- Minimal learning curve
- Strong community support

While tools like **IntelliJ** are excellent (and yes, I hear you, IntelliJ enthusiasts!), we'll defer discussions about premium IDE licenses until we've demonstrated concrete business value - typically after our first production deployment shows measurable ROI. This pragmatic approach helps us maintain focus on delivering value before optimizing developer preferences. This standardization is temporary - once we've established our foundation and demonstrated success, we can revisit IDE choices and accommodate team preferences.

And at this point, our **Neovim** users get a special hug from me - I fear they won't get it anywhere else - and I better stop here briefly to wipe away my tears.

1.4.2 Node.js

Ensure you have Node.js installed. If not, download it from <https://nodejs.org/>. Always use the latest LTS (Long Term Support) version.

At the time of writing this document, that's version **v22.11.0**.

Even better would be **installing a Node Version Manager** and switching between versions as needed¹⁰.

MacOS and **Linux Users**

For macOS or linux users, I recommend using **n** as your Node.js version manager. Install it globally via:

```
1 npm install -g n 
```

n provides seamless Node.js version management with minimal overhead. For comprehensive documentation and usage examples, refer to the official repository: <https://github.com/tj/n>

Windows Users

For Windows users, use **nvm-windows** - it's specifically designed for Windows environments. You can download it from <https://github.com/coreybutler/nvm-windows>

¹⁰ I can imagine that in larger enterprises, you can't access the internet from Windows Terminal and proxy servers need to be configured. There's probably a 2-page document on Confluence that's outdated and links to 10 other pages. In banks and insurance companies, getting this access can take up to a month (not joking). Until all approvals come through, we work with memory sticks (if they're not blocked too).

Corporate Reality

Only in huge enterprises (especially in Germany and Switzerland) you might find node-package in the internal software package system, but no `nvm-windows`^a.

If we can't install `nvm` on Windows Machines (and Azure AKS is still pending approval - which would allow us to work in containers), then everyone must install exactly the same version that also runs on the deployment server. And **just like that**, the architect has made themselves unpopular at the very beginning of the project^b.

^a Don't be surprised - it's not uncommon for developers to lack admin rights on their own machines in German speaking EU. In these cases, I usually recommend using your vacation bonus to buy an Apple notebook for development. I've probably already written a VPN driver for your company's Mac setup anyway. You can then use the Windows laptop for time tracking, reading emails, and responding to MS Teams messages. *BTW: How can anyone expect developers to work without admin rights on their own machines? I find that horrifying.*

^b But beware: If the company's internal cloud (like Azure) isn't configured with the latest LTS version (*I don't always have admin rights on internal cloud servers and need to ask the CTO to request one for me, assuming I don't feel like getting stoned to death*), we have to align with whatever version is there. It's not uncommon that I need to call India, South Korea, New Zealand, or Hungary to change my password. God help you if it's India - you'll definitely need a **phonetic alphabet chart**, and if the password contains special characters (*and it must*), you'll spend an afternoon on it and probably need to continue the next day. If we need to switch to an earlier LTS that doesn't support `root async/await`, I might even need to downgrade my ELVIS import generators. (ELVIS, our Nx Mono Repo, needs to be able to import all apps regardless of framework - more on that later).

2 Create a New Nx Workspace (Angular Setup)

This documentation is primarily designed for Angular experts¹¹ seeking a fast-track introduction to the NestJS ecosystem.

Impatient Minds

For immediate setup, jump to [Final Confirmation Chapter](#). For detailed explanations and rationale behind this setup, continue with the following chapters.

2.1 Give Your Workspace a Name

After all this talk about relationships and configurations, it's time to get hands-on and create something new. Let's bring our first **Nx workspace** to life - we'll call it **Elvis**, because it's about to shake things up! Open your terminal and run the following command:

```
› npx create-nx-workspace@latest elvis ↵ # or your preferred name
```

Déjà Vu

Well, well, well... After the very first command, it looks like our Nx buddies might have spent some time in the React.js world before. While in classic Angular we'd go with `npm i -g @angular/cli@latest` followed by `ng new <my-angular-app>`, our Nx team embraces the `npx` approach (suspiciously similar to `npx create-react-app`).

At this point, Junior Developers often ask if `npx` has something to do with `nx`, and when we installed `npx`. Let's clear this up:

What is npx?

`npx` is a package runner tool that comes bundled with `npm` (since version 5.2.0)^a. It's automatically installed when you install Node.js. Its main purpose is to execute `npm` package binaries without having to install them globally. Think of it as a way to “*try before you globally install*” - perfect for one-off commands like creating new projects.

^a This was a write-up on Medium by Kat Marchán, `npm` CLI team's green-haired digital dryad medium.com/@maybekatz/introducing-npx-an-npm-package-runner-55f7d4bd282b. My goodness, it's been 7 years - when did that happen?

¹¹ This chapter is dedicated to [@DanielGlejzner](#) and [@Armandotrue](#), two talented young colleagues who tirelessly share their Angular expertise with the community.

Does npx install packages locally?

No, npx doesn't install packages permanently. It works like this:

- First checks if the package is already installed locally
- If not, downloads the package to a temporary cache
- Executes it from there
- After execution, removes the downloaded files

So for commands like `npx create-nx-workspace`, the package itself isn't installed permanently - it's temporarily cached, used, and then cleaned up. However, the `workspace` it creates is permanent and includes all necessary dependencies in your project's `node_modules`.

2.2 Confirm Download & Workspace Installation Start

Alright, let's continue with the CLI interaction. At this point, **Nx** asks for permission to install the workspace:

```
Need to install the following packages:
create-nx-workspace@20.3.0
Ok to proceed? (y)  # press Enter
```

Makes perfect sense to answer with **(y)** here - just hit Enter and let's get this party started.

Why CLI Isn't a Lottery

Pro tip: Throughout the setup, use your **up/down arrow keys** to browse through the options and press **Enter** to confirm. And no, randomly mashing arrow keys while blindly hitting Enter isn't a valid selection strategy - though I've seen people try! Let's make our choices deliberately, shall we?

2.3 Choose Your Stack: Angular

```
NX Let's create a new workspace [https://nx.dev/getting-started/intro]
```

```
? Which stack do you want to use? ...
```

| | |
|----------|--|
| None: | Configures a TypeScript/JavaScript monorepo. |
| React: | Configures a React application with your framework of choice. |
| Vue: | Configures a Vue application with your framework of choice. |
| Angular: | Configures a Angular application with modern tooling. |
| Node: | Configures a Node API application with your framework of choice. |

What do you think, dear Angular experts, which option should we choose here? Angular, of course. Let's select it and continue by pressing Enter.

2.4 Multiple Projects Setup: Integrated Monorepo

```
✓ Which stack do you want to use? · angular
? Integrated monorepo, or standalone project? ...
Integrated Monorepo: Nx creates a monorepo that contains multiple projects.
Standalone:          Nx creates a single project and makes it fast.
```

Here we'll select "Integrated Monorepo" since we want to create multiple projects (apps) - at least one Angular app (which Nx will set up for us shortly) and later we'll add a Nest.js project (app). Press Enter to confirm.

2.5 Set Your Initial Application's Name

```
✓ Which stack do you want to use? · angular
✓ Integrated monorepo, or standalone project? · integrated
✓ Application name · fe-ng-base # or your preferred name
```

Next, Nx asks for the application name. Let's name it **fe-ng-base** (short for Frontend Angular Base) and press Enter (You can choose any meaningful name here).

2.6 Choose Your Bundler

```
✓ Which stack do you want to use? · angular
✓ Integrated monorepo, or standalone project? · integrated
✓ Application name · fe-ng-base
? Which bundler would you like to use? ...
esbuild [ https://esbuild.github.io/ ]
Webpack [ https://webpack.js.org/ ]
```

We'll answer the question "Which bundler would you like to use?" with "esbuild". Why? because it's faster. (Written in GoLang, optimised for JS/TS, lower memory footprint, native ESM support, well integrated with Nx,... etc. In the context of our monorepo, esbuild's speed and efficiency will become increasingly valuable as our project grows.)

Press Enter to confirm.

The Architect's Wisdom

While esbuild is very fast, it's worth noting that Rust-based **Rolldown** github.com/rolldown/rolldown is emerging as a potentially even faster alternative. For React apps, you would have **Vite**, **Webpack**, and **Rspack** as alternatives. While it would be possible to adapt Rolldown for Nx Angular apps with some effort, the community will likely implement this soon, so I'd wait for that. Currently, we know that Vite plans to replace both esbuild and Rollup with Rolldown. rolldown.rs/guide/#why-rolldown

2.7 Set Default Stylesheet Format

```
✓ Which stack do you want to use? · angular
✓ Integrated monorepo, or standalone project? · integrated
✓ Application name · elvis
✓ Which bundler would you like to use? · esbuild
? Default stylesheet format ...
CSS
SASS(.scss)      [ https://sass-lang.com   ]
LESS             [ https://lesscss.org     ]
```

At this point, we'll select **SASS(.scss)** as our default stylesheet format. This choice is strategic because SASS offers several key advantages for enterprise applications:

- Powerful nesting capabilities and variables
- Reusable mixins and functions
- Better code organization through partials
- Built-in color manipulation
- Advanced math operations

We'll later provide global SCSS mixins, accessibility features, and responsive design patterns out of the box (I'm planning to write a dedicated ELVIS book about these topics).

In an enterprise workspace, developers should focus on implementing critical features without getting bogged down by UI design, responsiveness, and accessibility concerns. However, these crucial and complex aspects need to be professionally implemented from the very beginning.

While product managers prioritize features for implementation, design should come from designers and usability experts who naturally maintain close contact with product management (domain experts, or “Fachleute, Fachabteilung” as we say in German-speaking EU). The architect serves as the crucial bridge between these groups.

A Wannabe Mathematician's Design Awakening

In the projects I support, I advocate for using **Figma** as our design platform. While I'm not a Figma expert myself I know enough about Figma to demonstrate its value to business departments and designers (*despite my background in Theoretical Computer Science with focus on Mathematics and later Media Informatics (UI/UX) - and yes, until that second degree, I thought I could design! Reality was quite the eye-opener*).

Often though, this isn't even necessary as they're already familiar with Figma. BTW, we've developed efficient methods to seamlessly transfer these designs into Angular applications.

This is where our SCSS mixins prove their true value. The automation capabilities that SCSS

offers are remarkable. I can't emphasize enough how underrated SCSS is - it doesn't receive the recognition it deserves in my opinion.

As the connecting element, the architect orchestrates the collaboration between product management (business departments), designers, and the development team.

Trust Me, I've Seen Things...

Allowing designers to work directly with developers, or letting developers handle design, would lead to disastrous results in 99% of cases^a.

^a Dear Niklaus, I'm thinking of you in this very moment...

Alright, let's continue with the CLI interaction. When prompted for the stylesheet format, select **SASS(.scss)** and press Enter.

2.8 SSR and SSG/Prerendering Setup

The next question appears:

Do you want to enable Server-Side Rendering (SSR) and Static Site Generation (SSG/Prerendering)?

```
✓ Which stack do you want to use? · angular
✓ Integrated monorepo, or standalone project? · integrated
✓ Application name · elvis
✓ Which bundler would you like to use? · esbuild
✓ Default stylesheet format · scss
? Do you want to enable (SSR) and (SSG/Prerendering)? ...
Yes
No
```

Let's take a brief detour to explore these two topics and then return to this point.

2.8.1 Server-Side Rendering (SSR)

First, let's clarify Server-Side Rendering (SSR)¹². For Angular experts reading this document - you're already familiar with all this. I'm just briefly recapping for completeness. Detailed information can be found on Angular's new website: angular.dev/guide/ssr

SSR is a technique where Angular renders the initial HTML on the server for each request. This HTML is then sent to the client along with the necessary JavaScript code. Once loaded, the JavaScript hydrates the application, adding interactivity to the pre-rendered HTML.

¹² While I haven't extensively tested the newest features yet I've been following the development closely. I'm planning to write a dedicated document/book about this topic with other Angular experts.

My journey with SSR has been interesting. In my early Angular implementations I couldn't optimize performance effectively and eventually abandoned it.

However, Angular now offers impressive improvements, particularly in Performance Metrics like **Time to First Byte (TTFB)**, thanks to optimizations in the rendering process.

Early Angular SSR hydration attempts were “destructive”. This meant that the loaded JavaScript code didn't properly register event handlers with already rendered DOM elements. Instead, the framework simply deleted all DOM elements and started rendering from scratch. This is why older Angular applications show a brief “flash” during initialisation where the already rendered page momentarily disappeared.

Regarding hydration, there has been a focus on making the hydration process **non-destructive**, meaning the client-side JavaScript doesn't need to rebuild the DOM but instead enhances it, improving initial load time and user experience. New hybrid rendering APIs have been introduced, allowing developers to mix SSR with client-side rendering for better performance where needed.

🐾 The Art of Web Hydration 🍲

Ever wondered why your grandma's old chicken soup recipe took hours to make, while instant soup is ready in minutes? Well, web hydration is kind of like that instant soup magic, but for websites!

Picture this: You're staring at a bowl of dried soup powder (that's your static HTML from the server). Sure, it looks like soup, but try eating it. Then add hot water, and suddenly you've got yourself a proper soup. That's exactly what happens with web hydration, minus the actual soup.

When a server sends your browser a webpage, it's like getting that packet of dried soup. It looks like a website, but try clicking those buttons and... nothing happens. That's because it's just a static snapshot, as interactive as a printed screenshot. But then JavaScript swoops in like hot water to our rescue, bringing everything to life!

Suddenly those lifeless buttons spring into action, forms start working, and your website transforms from a digital statue into a living, breathing application.

Could we explain Angular's approach to SSR to a web expert without Angular knowledge? Yes, we can. No matter which library or framework we use, they all need a root anchor, an id where everything begins. In a React.js world, it would look something like this (highly simplified), assuming we have:

```
1 <div id="root">
2   <!-- this is where the life begins -->
3 </div>
```

In Angular, it would look approximately like this (Angular makes it appear even more magical):

```
1 <app-root>
2   <!-- this is where the life begins -->
3 </app-root>
```

A static content (without SSR) inside could look like this (now general and not framework-specific):

```
1 <div id="root">
2   <div class="loading">Loading... </div>
3 </div>
```

And when dynamic parts (JavaScript, Data, ...) are loaded, the DOM is modified. Here, the loading part disappears and is replaced by another element, that need at least 2 DOM operations: Delete and Insert, afterwards we would have:

```
1 <!-- Then after JavaScript loads and executes: -->
2 <div id="root">
3   <div class="user-profile">
4     <h2>Welcome, Inspector Clouseau</h2>
5     <button class="likes-button">
6       Likes: 0.5
7     </button>
8   </div>
9 </div>
```

What's the difference with SSR? Through SSR, the client would receive something like this - a piece of HTML and some JavaScript code, and this Code will soon bring the already loaded parts to life:

```
1 <!-- Static Part: What to show -->
2 <div class="user-profile">
3   <h2>Welcome, Inspector Clouseau</h2>
4   <button class="likes-button">
5     Likes: 0
6   </button>
7 </div>
8 <!-- Dynamic Part: What to do -->
9 <script>
10 document
11   .querySelector(".likes-button")
12   .addEventListener("click", () => {
13     const currentLikes = parseInt(
14       this.textContent.split(":")[1],
15     );
16     this.textContent = `Likes: ${currentLikes + 1}`;
17     /**
18      * Real implementation might include:
19      * - API calls, State management,
20      * - Real-time updates (incl. User Interaction)
21      * - Error handling, Loading states
22      * - Analytics tracking, and much more
23      * */
24   });
25 </script>
```

This means there's a mechanism (called hydration) that (when activated) always expects two components: a static part and a dynamic part. The static part provides immediate visual content, while the dynamic part adds interactivity through JavaScript. Unlike client-side rendering, hydration avoids expensive DOM replacement operations - the view is already there and only needs to be 'brought to life' with event handlers and state management. This is the essence of web hydration.

Now some readers might wonder why such a seemingly simple requirement can be so enormously complex? The answer is simple: **Web Hydration** is a complex process that involves multiple layers of abstraction, intricate browser APIs, and a delicate dance between server and client. It's a bit like a magic trick - it looks simple on the surface, but behind the scenes, there's a lot of *sleight of hand* going on.

This complexity arises from several intertwined challenges which we can principally categorize into 3 groups:

- Component Lifecycle

- Initial HTML from server must match exactly what the client would render
- Event handlers need precise reattachment without rebuilding the DOM
- Component states must synchronise seamlessly
- User interactions during hydration need careful handling

- Route Management
 - Server must handle the same routes as the client
 - Deep links must work immediately
 - Client-side navigation needs to take over smoothly
 - Route guards and resolvers must function correctly
- State & Data Flow
 - Server state must transfer to client without loss
 - API calls need deduplication
 - Client cache needs proper initialization
 - Real-time updates must integrate smoothly

While the audience of this document consists of Angular experts who are already familiar with most of these details, we won't delve into the depths here (although that's where it gets more interesting). Instead, we'll focus on setting up an Nx workspace.

But at this point, we should briefly mention Static Site Generation (SSG/Prerendering), as our Nx CLI asked: 'Do you want to enable Server-Side Rendering (SSR) and Static Site Generation (SSG/Prerendering)?' We've covered SSR, now let's look at SSG.

2.8.2 Static Site Generation (SSG/Prerendering)

SSG is like SSR's efficient cousin - instead of rendering pages on each request, it pre-renders them during build time. Imagine you're running an e-commerce site:

- SSR: Renders pages when users request them
- SSG: Renders all pages once during deployment

Key differences:

- Build vs Runtime: SSG generates HTML files during build, not per request
- Performance: SSG serves pre-built static files (extremely fast)
- Resource Usage: No server computation needed for each visit
- Perfect for: Documentation, blogs, landing pages, product catalogs
- Less ideal for: Highly dynamic content, real-time data

Think of SSG as a printing press (print once, distribute many) while SSR is more like a real-time translator (translates as needed). In Angular's implementation, you can use both strategies where they make the most sense.

2.8.3 Challenges

For Angular developers new to **SSR** and **SSG**, the journey involves several important adaptations to their development approach.

The first noticeable change comes with **build** and **deployment** processes. Instead of a single browser-targeted build, you'll work with multiple build outputs targeting different platforms. Your application needs to run both in the browser and on a Node.js server, which introduces new complexity to the deployment pipeline.

Code architecture requires a more thoughtful approach. Things that were once taken for

granted, like direct access to browser APIs (window, document, localStorage), now need careful consideration. Your code must gracefully handle **both server and browser environments**, often requiring **platform-specific checks and conditional logic**.

Data handling becomes more sophisticated. You'll need to manage the flow of data between server and client, handle API calls appropriately in different environments, and implement effective caching strategies. The challenge lies in maintaining a **consistent state** across these different **execution contexts**.

The **development workflow** itself evolves. **Builds** take longer as they target **multiple platforms**. **Debugging** becomes more complex as issues might manifest differently on server versus client. **Testing** needs to account for both environments, and error handling must be environment-aware.

Common challenges include **avoiding infinite loops in lifecycle hooks**, **preventing memory leaks on the server**, and **handling browser-only third-party libraries**. These aren't insurmountable obstacles, but they require a broader understanding of the **application's execution context**.

The key is understanding that your application now lives in two worlds - **server** and **client** - and needs to transition smoothly between them. This dual-nature brings powerful benefits but requires a more nuanced development approach.

2.8.4 Architect's Notes

When do architects change their proven principles?

Most architects have a portfolio of diverse requirements and choose their proven tools according to the situation and goals. For me, **NestJS** is a proven tool, and I've had exclusively good experiences with it. However, even NestJS must sometimes make way for alternatives.

Consider scenarios where **raw performance** becomes critical - here, switching to **Go** or **Rust** can provide significantly better runtime performance. Enterprise environments often mandate **Java** for better integration with existing systems, especially when API teams develop exclusively in Java and need to maintain technological consistency across their stack.

The decision to switch can also be driven by specific technical demands: **microservices** requiring extreme **efficiency in resource usage**, platform requirements demanding **native capabilities**, or when **SSR/SSG** mechanisms simply offer more elegant solutions for particular use cases. **Real-time processing** with **strict latency requirements** might also necessitate a different technological approach.

These examples illustrate a broader principle: while we should value our proven tools, we must remain flexible enough to embrace better solutions when they **truly serve our specific needs**. The key lies not in changing technologies for change's sake, but in recognizing when a different approach genuinely better serves our architecture's goals.

When it comes to technology choices, I'm always ready to pivot when **quantum leaps in performance** or **complexity management** emerge. It's like upgrading from a horse and buggy to a Tesla - sometimes the jump is just that dramatic. However, the reverse can also be true - sometimes you **might trade in your Tesla** because what you actually need right now are some cows, sheep, or chickens. It's about **choosing the right tool for the current need**, not always about having the most advanced technology.

Our development landscape has seen **numerous transformative shifts** that reshaped how we work. **JQuery** gracefully bowed out as native **DOM APIs** evolved, while **SOAP** messages transformed into sleek **REST APIs** before **GraphQL** emerged to change the game again.

We witnessed the shift from **monolithic PHP applications** to modern microservices, and saw the evolution **from Angular.js to Angular 2+** - a revolution disguised as an upgrade. Even **Flash**, once ubiquitous, gracefully retired in favor of **HTML5**.

Today's quantum leaps are equally compelling: **Rust** offers tenfold performance improvements for **CPU-intensive tasks**, while **Go** reduces **memory footprint** by orders of magnitude in **high-concurrency scenarios**. **Bun** demonstrates significantly faster npm install times compared to **Node.js**, and **edge computing platforms** achieve **microsecond responses** versus traditional cloud deployments.

Meanwhile, the rapid evolution of **AI**, from **narrow AI** to the pursuit of **AGI** (Artificial General Intelligence) and theoretical **ASI** (Artificial Superintelligence), is reshaping how we think about software architecture and system design. The integration of AI capabilities has become **a game-changer**, transforming traditional programming paradigms and opening new possibilities in **automated coding, optimization, and decision-making processes**.

Sometimes, the factors that drive these changes **extend beyond pure performance metrics**. We're seeing fundamental shifts in how systems are designed, developed, and maintained. While narrow AI is already enhancing our development workflows through tools like copilots and code analyzers, the potential impact of more advanced AI systems could revolutionize our entire approach to software architecture.

🕒 **PHP: Permanently Heartwarming Platform**

PHP has shown remarkable resilience, initially thanks to frameworks like Yii, Symfony, Laravel, and WordPress, and later through its community's incredible adaptability. It remains the go-to language for catalogs, blogs, and static CMS systems. PHP has become something of the COBOL of our time - steadfastly reliable and surprisingly enduring.

Even its name evolution tells a story of growth - from "Personal Home Page" to "PHP: Hypertext Preprocessor", following the clever recursive naming pattern popularized by GNU ("GNU's Not Unix"). Like many developers, I built my first websites with PHP and financed my university studies through PHP projects. It's fascinating - I've yet to meet anyone who has worked with PHP and doesn't maintain a fond appreciation for it.

This staying power demonstrates an important principle in technology: sometimes the most practical solution, even if not the most elegant, becomes a lasting one. PHP's pragmatic approach to web development has earned it a permanent place in our technological landscape.

That's why I simply had to write something for PHP. And now I need a tissue to wipe away my tears.

2.8.5 SSR & SSG in Nx: Decision Time

The focus of this document is Nest.js for Angular experts. We've made a conscious decision to maintain a clear separation: Angular for the client, Nest.js for the server. (While we could explore combining these technologies - and there's nothing technically preventing us from doing so - that discussion would exceed the scope of this document. We'll save that fascinating topic for the upcoming ELVIS book.)

```
✓ Which stack do you want to use? · angular
✓ Integrated monorepo, or standalone project? · integrated
✓ Application name · elvis
✓ Which bundler would you like to use? · esbuild
✓ Default stylesheet format · scss
? Do you want to enable (SSR) and (SSG/Prerendering)? ...
Yes
No
```

Please select **No** here and press Enter.

2.9 Choose Your Test Runner (E2E)

```
✓ Which stack do you want to use? · angular
✓ Integrated monorepo, or standalone project? · integrated
✓ Application name · elvis
✓ Which bundler would you like to use? · esbuild
✓ Default stylesheet format · scss
✓ Do you want to enable (SSR) and (SSG/Prerendering)? · No
? Test runner to use for end to end (E2E) tests ...
Playwright [ https://playwright.dev/ ]
Cypress [ https://www.cypress.io/ ]
None
```

Testing is an extraordinarily critical subject - one that could easily warrant its own book. While this topic isn't strictly part of this document, I'd like to take this opportunity to highlight some **key principles**.

2.9.1 Key Principles

You have complete flexibility in choosing any of these options for your E2E testing needs.

While I lean toward **Playwright** for **Azure Cloud** environments, **Cypress** shines brilliantly in its own right. Its intuitive interface is a masterpiece of design, offering an exceptional developer experience. The selector playground is particularly impressive - it makes element selection effortless and speeds up test creation significantly.

Corporate Policies Matter

As enterprise architects and developers, we understand a fundamental reality: tools must be pre-approved before they can be implemented. Having introduced numerous new tools, languages, and frameworks into traditionally conservative enterprise environments, I'm intimately familiar **with the challenges** this process presents.

When starting a new project, if the CTO indicates that technology X is off-limits and I know of a viable alternative Y, I'll pragmatically choose Y. However, when X is truly indispensable, we must strategically make our case.

This involves creating a **comprehensive**, yet **accessible comparison** of advantages and disadvantages, **presented visually** in a way that **resonates with board-level executives**.

Success in these situations requires more than technical merit - it demands clear communication of business value in non-technical terms. While **there's never a 100% guarantee of approval**, a well-structured, business-focused presentation significantly improves our chances of gaining support for essential technical innovations.

Remember: The key is choosing our battles wisely and presenting our case in a language that aligns with business objectives.

Newcomers to **E2E** testing often find **Cypress's learning curve** remarkably gentle, thanks to its **well-thought-out design and outstanding documentation**. The vibrant, knowledgeable **community** consistently provides top-tier support and valuable insights.

For Azure Cloud environments, I prefer Playwright due to its compatibility with **enterprise policies** regarding **binary downloads and pipeline permissions**. Its broader feature set proves particularly valuable in complex enterprise scenarios. (While obtaining special approvals for Cypress is possible, the process can be time-consuming and potentially delay project timelines.)

As an aside, **Puppeteer** would be a wonderful addition to this toolkit, especially for **AI development** where its capabilities extend beyond traditional testing. (A note to our **Nx Core Team** - considering Puppeteer would be appreciated.)

2.9.2 Challenges

Even brilliant architects can find themselves in challenging or seemingly absurd situations when navigating German-speaking EU enterprise companies without prior experience. The fundamental truth remains: software isn't properly validated without thorough testing.

Consider this scenario: Your team has already invested significant effort creating comprehensive E2E tests with Cypress. Then you discover that your CI/CD pipelines won't support Cypress binaries, and the tool hasn't received corporate approval. The consequence? All tests need to be rewritten - a classic own goal for the architect. Developers will understandably be frustrated, having dedicated substantial time to crafting excellent E2E tests. (*And if we're forced to use Selenium, that's a frustration I personally share*)

This creates a complex dilemma:

- If you delay E2E testing while waiting for framework approval, developers might question your architectural judgment
- If you implement tests that later need complete rewriting, you waste valuable development resources

The solution lies in proactive communication and planning:

- Inform the development team upfront about the approval process and potential tooling constraints
- Engage with product management early to reserve time and budget for potential testing implementation or migration
- Account for framework approval timelines (which could take 1 week to 3+ months) in the project planning
- Set aside dedicated sprints (minimum 2) for implementing E2E tests once tools are approved

This approach prevents surprises and conflicts with both development teams and product owners who might otherwise push back against “unexpected” testing efforts later in the project lifecycle.

Remember: In enterprise environments, technical decisions must align with both organizational policies and project timelines from the very beginning.

2.9.3 Decision Time

Choose the tool that best fits your specific context - both options have their strengths, and either could be the perfect choice depending on your needs.

```
? Test runner to use for end to end (E2E) tests ...
Playwright [ https://playwright.dev/ ]
Cypress [ https://www.cypress.io/ ]
None
```

Make your choice and press Enter.

2.10 CI Provider Section

```
✓ Which stack do you want to use? · angular
✓ Integrated monorepo, or standalone project? · integrated
✓ Application name · elvis
✓ Which bundler would you like to use? · esbuild
✓ Default stylesheet format · scss
✓ Do you want to enable (SSR) and (SSG/Prerendering)? · No
✓ Test runner to use for end to end (E2E) tests · playwright
? Which CI provider would you like to use? ...
```

```
GitHub Actions
Gitlab
Azure DevOps
BitBucket Pipelines
Circle CI
```

```
Do it later
```

Remote caching, task distribution and test splitting are provided by Nx Cloud. Read
↪ more at <https://nx.dev/ci>

While CI/CD configuration is a particularly engaging aspect of development (and one I personally enjoy), we'll select "Do it later" for this setup. This topic deserves comprehensive coverage and will be addressed in detail in the upcoming ELVIS book.

Select **Do it later** and press Enter to proceed.

2.11 Remote Caching

```
✓ Which stack do you want to use? · angular
✓ Integrated monorepo, or standalone project? · integrated
✓ Application name · elvis
✓ Which bundler would you like to use? · esbuild
✓ Default stylesheet format · scss
✓ Do you want to enable (SSR) and (SSG/Prerendering)? · No
✓ Test runner to use for end to end (E2E) tests · playwright
✓ Which CI provider would you like to use? · skip
? Would you like remote caching to make your build faster? ...
(can be disabled any time)
Yes
No - I would not like remote caching
Read more about remote caching at https://nx.dev/ci/features/remote-cache
```

While remote caching is an excellent feature that significantly improves build performance, we'll select "No" here for a critical reason: Enterprise EU corporations typically have **strict data protection policies** that prevent using external cloud services for build artifacts.

The rationale behind this limitation:

- **EU data protection regulations (GDPR)** require careful handling of any data leaving corporate networks. Build artifacts might contain **sensitive information** or **intellectual property** and Internal **audit requirements** often mandate that all build processes remain within controlled environments.
- Many enterprises require that development artifacts stay within their own infrastructure and implement their own internal caching solutions that comply with their security policies.

Select **No** and press Enter to continue with the setup.

🐾 Order 66: When Caches Turn to the Dark Side

Accidentally choosing “Yes” to remote caching:

While technically possible to connect to Nx Cloud from your machine when working remotely (with VPN conveniently “disabled”), this might trigger some... interesting corporate responses:

Scenario A - Working From Home

At 4:30 AM, your peaceful sleep might be interrupted by a tactical security team, complete with dramatic blue lights illuminating your neighborhood. You’ll have the unique experience of being escorted in full restraints while your neighbors get an early morning light show.

Scenario B - At the Office

A security team member, armed with what appears to be a butterfly net, will swiftly apprehend you. You’ll be expertly guided to the exit, while your laptop - now considered a potential security threat - is whisked away for “further investigation.” **Remember:** In enterprise security, there’s no such thing as “it’s just a build cache.”

2.12 Final Confirmation

- ✓ Which stack do you want to use? · angular
- ✓ Integrated monorepo, or standalone project? · integrated
- ✓ Application name · elvis
- ✓ Which bundler would you like to use? · esbuild
- ✓ Default stylesheet format · scss
- ✓ Do you want to enable (SSR) and (SSG/Prerendering)? · No
- ✓ Test runner to use for end to end (E2E) tests · playwright
- ✓ Which CI provider would you like to use? · skip
- ✓ Would you like remote caching to make your build faster? · No

NX Creating your v20.3.0 workspace.

- ✓ Installing dependencies with npm
- ✓ Successfully created the workspace: elvis.

NX Welcome to the Nx community! 🙌

- 🌟 Star Nx on GitHub: <https://github.com/nrwl/nx>
- 🇺🇸 Stay up to date on X: <https://x.com/nxdevtools>
- 💬 Discuss Nx on Discord: <https://go.nx.dev/community>

3 Exploring the Nx Angular Workspace

3.1 Open Your Workspace in VS Code

Let's switch to our Nx workspace (in our case, it's named elvis) and launch VSCode from the terminal. For those who haven't set up the command line shortcut yet, you can refer to the "Install Recommended Software" chapter.

```
1 cd elvis ↵
```

and then start VSCode from the terminal, notice the dot (.) at the end of the command:

```
1 code . ↵
```

VSCode should open now.

3.1.1 VSCode UI Interface Layout and Terminology

Please click the **Explorer** icon in the **Activity Bar** (if it's not already selected) to view the workspace structure.




Figure 6: VSCode Explorer Icon & Other Primary Side Bar Icons

Refer to the VSCode chapter in this document for the interface terminology or check the official docs here: code.visualstudio.com/docs/getstarted/userinterface

I want to explain why I insist on using consistent terminology - specifically the one defined by the original source. (If I ever deviate, I'll explicitly mention why, which happens extremely rarely).

While discussing the **Explorer View** in VSCode might *seem trivial* now, besides the standard icons (*Search*, *Source Control*, *Run and Debug*, *Extensions*) on the **Activity Bar**, there are additional ones added by extensions, several of which we'll be installing.

When you're conducting a demo and ask colleagues to click on a specific icon, everyone needs to know the terminology and UI interface. Otherwise, they might miss an icon right in front

of them.

🕒 When Panic Meets Precision: A Tale of Remote Debugging

Consider a scenario where a junior programmer needs to fix a bug urgently (because a Product Owner bypassed the Scrum Master, sprint planning, and architect, choosing a junior developer as the target), and they call you for remote pair programming. If they're panicking, moving the mouse erratically, and unfamiliar with the terminology, you'll understand my point.

I've encountered situations where I had to ask colleagues to stop moving their mouse. I'd ask their monitor size and resolution, then have them take a ruler and move the mouse 12cm left and 7cm up to click the correct icon. Knowing the icon's name and its location bar makes everything much simpler.

3.1.2 Checkpoint: 001-elvis-after-creation

📄 Git Ready: ELVIS Has Entered the Repository

I have already created a repository for ELVIS and save each significant step we take together in a separate branch. This allows us to compare our current state with the repository at any time. The repository can be found here: github.com/nxpatterns/elvis

Before we examine the workspace content in detail, let's install some important extensions. But first, let's check our initial repository state. At this point, we should have the state of:

```
github.com/nxpatterns/elvis/tree/001-elvis-after-creation
```

3.2 Install Recommended VSCode Extensions

3.2.1 Create .vscode/extensions.json File

With VS Code open and the Explorer selected in the Activity Bar, we'll now install the recommended extensions. The most efficient approach is to create an `extensions.json` file in the `.vscode` folder at:

```
├── .vscode
│   └── {} extensions.json
```

Let's open the terminal in VS Code via **View > Terminal > New Terminal**. We'll set up a shortcut for this later, after installing our extensions.

For macOS/Linux, create the file with:

```
1 mkdir -p .vscode && touch .vscode/extensions.json ↵
```

3.2.1.1 Windows: Creating Dot-Folders

Creating folders beginning with a dot can be tricky in Windows. Either use **Git Bash**, **WSL**, or create the folder via **VS Code's Explorer** by right-clicking and selecting **New Folder**. Alternatively, use the command prompt with:

```
1 md ".vscode"
```

WSL (Windows Subsystem for Linux)

For Windows users, I strongly recommend WSL - a native Windows feature that provides a genuine Linux environment without the overhead of a virtual machine. Find detailed instructions here:

learn.microsoft.com/en-us/windows/wsl/install

Windows Alternatives When WSL Isn't an Option

Sometimes corporate policies prevent WSL installation. In such cases, here are two excellent alternatives:

Git Bash: A lightweight shell providing a Unix-like environment on Windows. It comes bundled with Git. If you've installed Git, you already have Git Bash. During installation, ensure the "Use Git from the Windows Command Prompt" option (or similar) is selected. Installation details may vary. More information here:

— git-scm.com/downloads

Chocolatey with Cygwin: My preferred option when WSL isn't available. Cygwin provides an excellent Unix-like environment on Windows, while Chocolatey, a Windows package manager, simplifies its installation tremendously. Learn more here:

— chocolatey.org and
— cygwin.com

3.2.2 Add Recommendations to `extensions.json`

Open the file and add these recommendations:

```
1 {  
2   "recommendations": [  
3     "dbaeumer.vscode-eslint",
```

```
4    "EditorConfig.EditorConfig",
5    "esbenp.prettier-vscode",
6    "mrm1nc.vscode-scss",
7    "nrwl.angular-console",
8    "redhat.vscode-xml",
9    "redhat.vscode-yaml",
10   "shd101wyy.markdown-preview-enhanced",
11 ]
12 }
```

Managing Existing Extensions

Since we've just created the workspace with Nx/CLI, there shouldn't be an `extensions.json` file yet. If one exists (you've likely already added your preferred extensions), simply merge these entries with your current recommendations.

Pro tip: VSCode will automatically merge and deduplicate entries if you paste them into an existing file. Just ensure proper JSON syntax is maintained.

3.2.3 Restart VS Code

Now, please **close** Visual Studio Code **entirely** and **reopen** it. A similar window will appear in the bottom right asking if you want to install the recommended extensions:




Figure 7: VSCode Install Recommended Extensions

Click **Install** as we'll need all of these extensions. (By the way, this document was created using the `shd101wyy.markdown-preview-enhanced` extension.)"

After installation, you'll spot a new addition to our **Activity Bar**: the **Nx Console** icon:



Figure 8: New Addition: Nx Console in VS Code Activity Bar

3.2.3.1 Refreshing Window without Restart

By the way, we can also refresh our current window without the need for a complete restart. However, this isn't always sufficient, and in some cases, a full close and reopen may still be

necessary¹³:

This can be accomplished through the **Command Palette** (You may recall this from our earlier chapter when we installed VS Code).

If not mapped differently, the **F1** key works across Windows, macOS and Linux. If not, then use the following shortcuts:

- Windows/Linux: **Ctrl+Shift+P**
- macOS: **(Cmd+Shift+P)**

Then type: Developer: Reload You should see in the **Title Bar** of the VSCode window a dropdown menu; select **Developer: Reload Window** and press **Enter**.



Figure 9: Developer: Reload Window

3.2.4 Checkpoint: 002-elvis-after-recommended-extensions

At this point, we should have the state of:

github.com/nxpatterns/elvis/tree/002-elvis-after-recommended-extensions

3.3 Angular Setup Review

Let's take a quick peek under the hood of our Nx workspace. Think of it as a pre-flight check – we want to make sure our Angular engines are purring smoothly before we add NestJS to our tech stack.

We'll kick the tires, run some tests, and make sure everything's ready for takeoff. Better safe than sorry, right? Nobody wants their codebase doing unexpected barrel rolls once we add NestJS to the mix!

3.3.1 Serve the Angular Application

Let's start by serving our Angular application. We'll use the Nx Console extension to do this. Click on the **Nx Console** icon in the **Activity Bar** to open the Nx Console.

¹³ There's a shortcut available (Command+R on macOS), but it's tied to a condition (isDevelopment). We'll modify this later to enable context-free reloading. (This will be covered later, at the latest in the upcoming ELVIS book.)

Then go to **Projects** → Project Name (in our case **fe-base**) → **serve** → **development** and click “Execute task” (Play Button). This will initiate a terminal session and start the Angular development server.

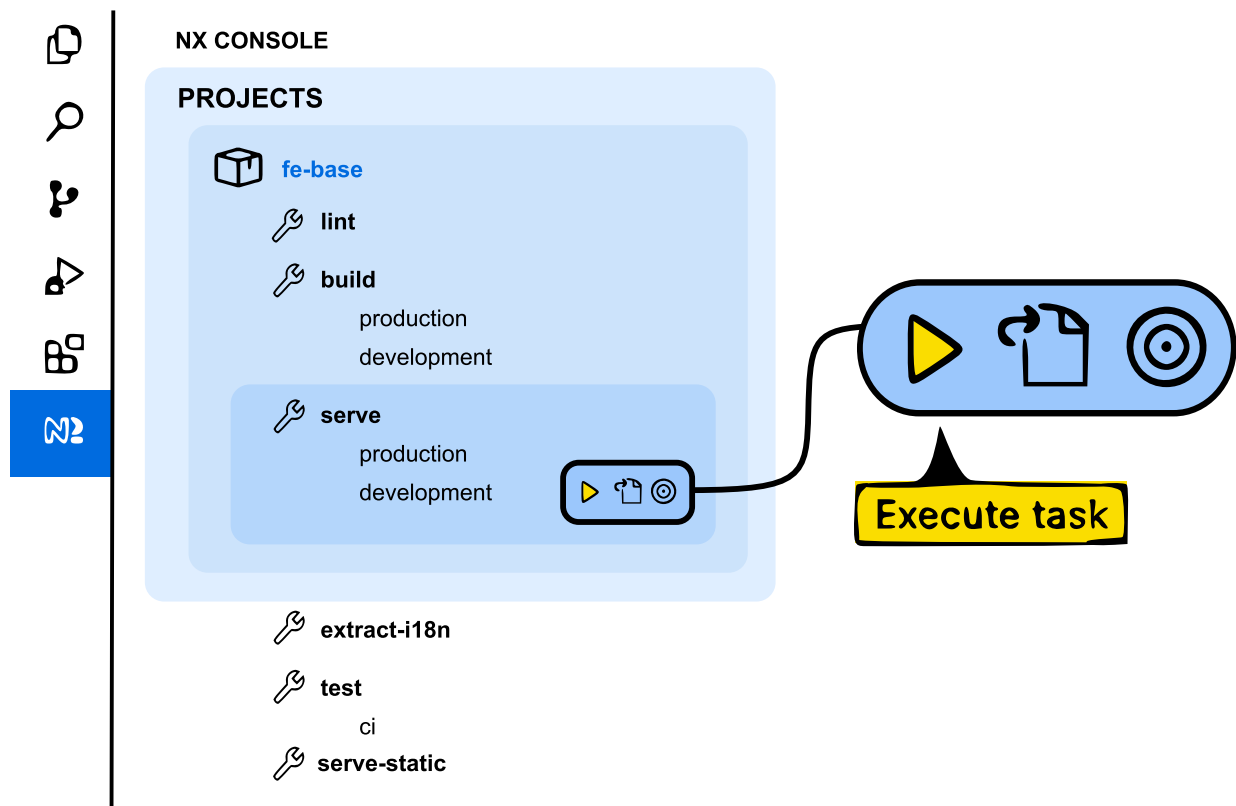


Figure 10: Nx Console Extension: Serve Angular App

For completeness, let’s briefly mention the functions of the other icons. (Details will be covered in the upcoming ELVIS book):



Figure 11: Other Icons Adjacent to “Execute Task”

i Streamline Your Development Workflow

Now, let’s make a mental note of this command, as we’ll soon edit our **project.json** to create a serve script. We’ll use this script to automate the process and assign a shortcut, enabling quick keyboard access in the future. This will be rather handy, as it’s a repetitive task we’ll be using daily throughout the project’s lifecycle - and who doesn’t love saving precious time?

Now let’s examine in detail which terminal commands are executed and their outcomes:

```
1 * Executing task: npx nx run fe-base:serve --configuration=development
2
3
4 > nx run fe-base:serve:development
5
6 Initial chunk files | Names          | Raw size
7 polyfills.js       | polyfills      | 90.20 kB |
8 main.js            | main           | 21.96 kB |
9 styles.css         | styles         | 109 bytes |
10
11                   | Initial total | 112.27 kB
12
13 Application bundle generation complete. [1.169 seconds]
14
15 Watch mode enabled. Watching for file changes...
16 NOTE: Raw file sizes do not reflect development server per-request transformations.
17 → Local:   http://localhost:4200/
18 → press h + enter to show help
```

Let's focus on these two lines first:

```
1 npx nx run fe-base:serve --configuration=development
2 nx run fe-base:serve:development
```

We've mentioned `npx` before; when `nx` (Nx/CLI) is locally installed (in our case, it's already installed as a dev-dependency in `package.json`), the installed version is called. If not (which can only happen if someone manually set up an Nx Workspace without CLI - rare but possible), it's downloaded, executed **but not installed**.

These subtle details **become crucial in cloud deployment environments**. Either `nx` must be installed **locally** as a **development dependency**, or **globally installed in the deployment pipeline before building** the app.

To avoid reinventing the wheel and debating naming conventions, I typically accept Nx/CLI's suggestion when creating an NPM scripts.

This means using `"fe-base:serve:development"` as the key (NPM script) and `"npx nx run fe-base:serve --configuration=development"` as the value, while **double-checking** that the correct version of `nx` (compatible with the used apps) is included as a dev-dependency in `package.json` ¹⁴.

Click the **Explorer** icon in the **Activity Bar** (if you've read this book/document in chapter order, you'll know exactly what the "Activity Bar" is and its location), scroll to the bottom and click **package.json** to add this small but essential entry:

¹⁴ This becomes particularly important when someone clones the project five years down the line and wants to run it locally.


```
1 // package.json (excerpt)
2 "scripts": {
3   "fe-base:serve:development": "npx nx run fe-base:serve --configuration=development
   ↪   --host=0.0.0.0"
4 },
```

You’ve likely spotted our sneaky addition: `--host=0.0.0.0`. The `--host` parameter is quite clever - it tells our application to listen for connections on all network interfaces (0.0.0.0), rather than just `localhost`.

This means your application becomes accessible from other devices on your network, which is jolly useful for testing on different devices (e.g. mobile phones etc.) or when running in containers.

In the `package.json`, `npx` isn’t strictly necessary since `nx` is locally installed. However, I include it because I’ve seen colleagues attempt to copy-paste these commands directly into deployment pipelines (*without having Nx CLI installed either globally or locally - yes, some Cloud experts actually manage to pull this off*). Including `npx` increases the chances of the pipeline finding and installing the correct package.

When npx Meets Bureaucracy

Fun fact: Everything’s a trade-off! Sometimes omitting `npx` makes more sense to quickly discover a missing package.

Picture this: You’re dealing with a heavily restricted pipeline, and your company’s artifactory only offers a prehistoric version of Nx. Yes, in some corporations, you need to beg the Cloud Team, write tickets, get CTO approval just to add an updated NPM package to the corporate repo.

By the time Security and Data Protection departments have finished their thorough investigation (approximately 6 weeks later), you might have grown a full beard. Don’t worry though - AI won’t be replacing this bureaucratic masterpiece anytime soon!

While making these changes, our Nx Serve task was still running in the terminal. Click on the terminal window and press CTRL + C to stop the task. You can then close this terminal tab by pressing any key.

Please note: Apart from the rightmost icons (`maximize` and `close`) in the `Panel-Tabs-Area`, all others action icons are context-sensitive.

In VSCode terminology, this entire `icon section` is called the `Actions Container`. The complete bar housing the `Panel Tabs` and these action icons is termed the “`Composite Bar`”. This terminology becomes increasingly crucial during remote pair programming sessions. I strongly recommend familiarising yourself with these terms from the outset.

In case the terminal window has closed (perhaps it was the only tab running the Nx Serve

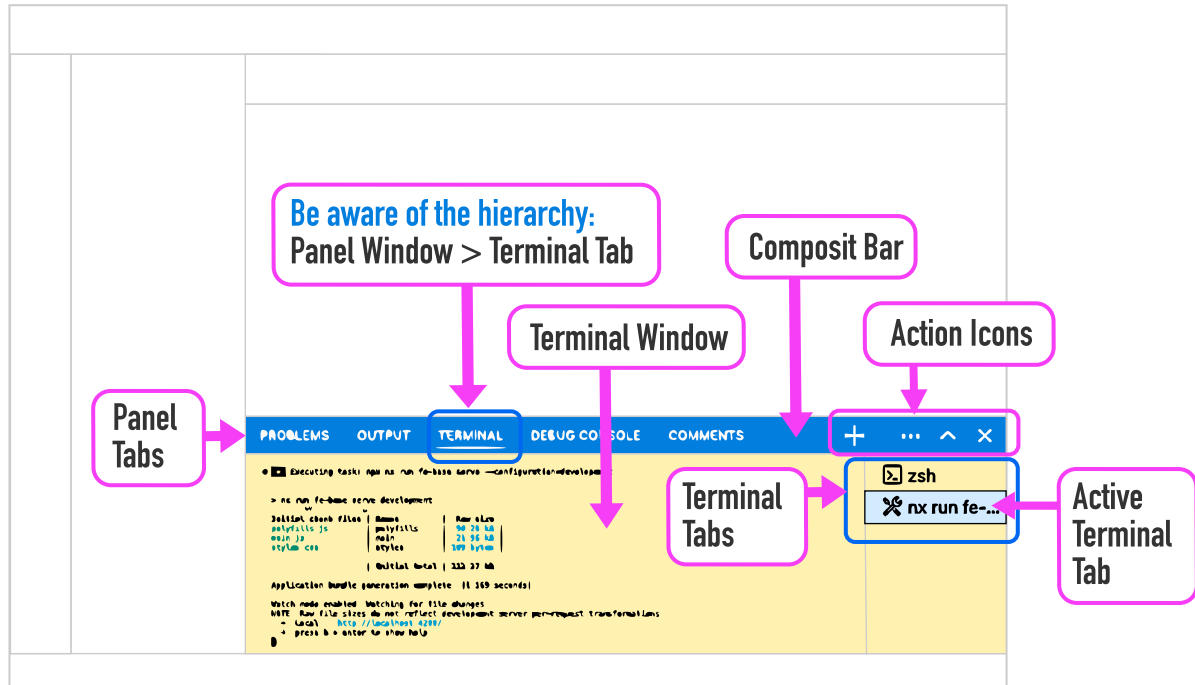


Figure 12: Terminal window: Access via “Terminal” in Panel Tab



Figure 13: Maximize and Close Icons

task), reopen it via Menu → Terminal → New Terminal¹⁵.



Figure 14: Reopen Terminal in VS Code

Let's test our newly added npm script. But first, a rather crucial point about our **naming convention**. We'll consistently use **this pattern** across all applications:

```
npm run <app-name> : <target> : <environment>
```

App names typically follow company conventions. I prefer adding prefixes like **fe-**, **be-**, or **mw-**, using lowercase with hyphens only. (This will be crucial for cloud pipeline configuration).

Common targets include e.g. **serve**, **build**, **test**, **lint**, **e2e**, **deploy**,..., while environments are typically e.g. **development**, **test**, **presentation**, **acceptance**, **production**,... The architect defines these with Product Management and Cloud teams.

In the Terminal, type this command and press Enter:

```
1 npm run fe-base:serve:development ↵
```

Here you'll notice **your machine's IP address** alongside **localhost**. This will prove rather handy later:

```
1 # ...
2 → Local:   http://localhost:4200/
3 → Network: http://<your-machine-ip>:4200/
4 → press h + enter to show help
```

Open your preferred browser, type **http://localhost:4200** in the address bar, and press Enter. (We'll later establish mechanisms to start the project, launch the browser, and navigate to the homepage with a single shortcut).

You should see your new Angular app with Nx Setup running smoothly in the browser (**Note:** *logos and images have been simplified to save printer ink - actual brand colors and design details in your browser will be more vibrant*):

3.3.2 Checkpoint: 003-serve-the-angular-application

At this point, we should have the state of:

¹⁵ Later, we'll establish unified shortcuts across platforms (macOS, Windows, Linux), store them in the workspace configuration, and maintain them in our technical reference guide.

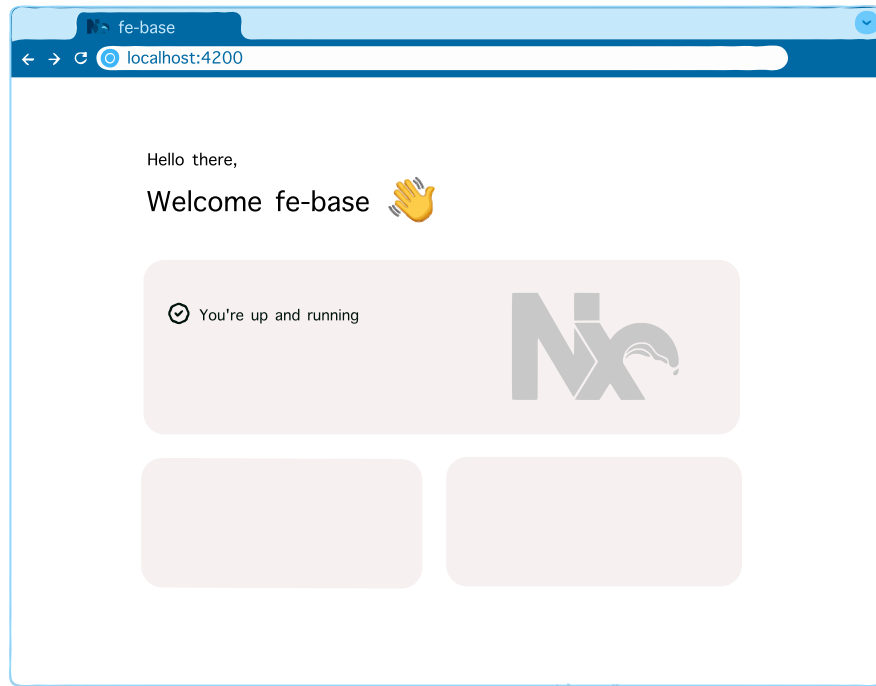


Figure 15: Nx Workspace - Angular Setup: Generated Welcome Page

github.com/nxpatterns/elvis/tree/003-serve-the-angular-application

3.3.3 Setting Up a Clean Workspace

Next, let's remove the Nx welcome page and update our `README.md`. (Please note the naming convention: “`README.md`” in uppercase letters for the filename only is used for the **root-level readme** (another name for **global readme**), all others are written as “`ReadMe.md`”. More on this later). Let's delete the file `apps/fe-base/src/app/nx-welcome.component.ts` and remove its import in `apps/fe-base/src/app/app.component.ts`:

```
1 import { Component } from '@angular/core';
2 import { RouterModule } from '@angular/router';
3 import { NxWelcomeComponent } from './nx-welcome.component'; // <-- Remove this line
4
5 @Component({
6   imports: [NxWelcomeComponent, RouterModule], // <-- Remove NxWelcomeComponent
7   selector: 'app-root',
8   templateUrl: './app.component.html',
9   styleUrls: ['./app.component.scss'],
10 })
11 export class AppComponent {
12   title = 'fe-base';
13 }
```

Remove line 3 with the `NxWelcomeComponent` import. Remove `NxWelcomeComponent` from the imports array in the sixth line. Let's review the complete file structure after removing those entries:

```
1 import { Component } from '@angular/core';
2 import { RouterModule } from '@angular/router';
3
4 @Component({
5   imports: [RouterModule],
6   selector: 'app-root',
7   templateUrl: './app.component.html',
8   styleUrls: ['./app.component.scss'],
9 })
10 export class AppComponent {
11   title = 'fe-base';
12 }
```

Our work regarding the `NxWelcomeComponent` is not yet complete. Please remove the `NxWelcomeComponent` selector `app-nx-welcome` from `apps/fe-base/src/app/app.component.html`, leaving only:

```
1 <router-outlet></router-outlet>
```

Now we will continue by editing the global `README.md`. Delete all content and replace it with:

```
1 # Elvis
2
3 A fresh take on Nx workspaces: Enterprise Level Visionary Innovation Stack.
```

Now let's test our VSCode extension "Markdown Preview Enhanced". If you haven't installed it yet, remember it was configured in `.vscode/extensions.json` under the `recommendations` section - if needed, please revisit the VSCode extensions installation chapter¹⁶.

Click the "Open Preview to the Side" icon in the editor's top-right corner (looks like an open book with magnifying glass) or use the keyboard shortcut:

- Windows/Linux: `Ctrl+K V`
- macOS: `Cmd+K V`

This will show you a live preview of your Markdown formatting:

The Markdown Preview Enhanced extension will be fundamental to our documentation strategy. We'll document the entire workspace architecture directly in the `README.md` file, includ-

¹⁶ marketplace.visualstudio.com/items?itemName=shd101wyy.markdown-preview-enhanced

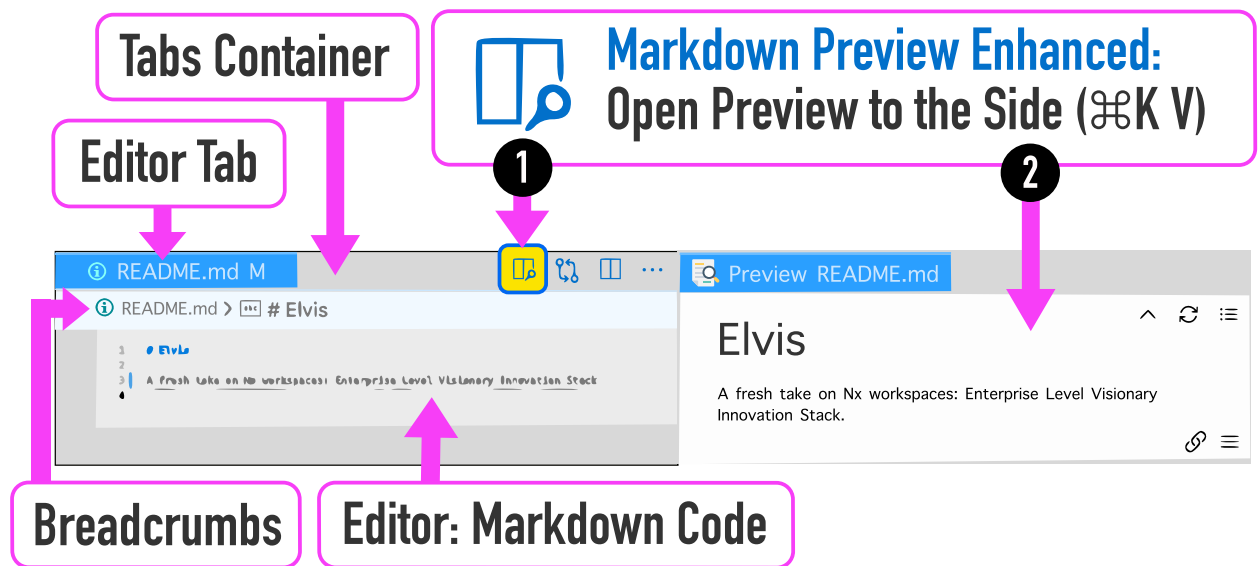


Figure 16: Markdown Preview Enhanced: Activate Live Preview

ing all architectural diagrams. By keeping the documentation in the repository and versioning it alongside the code, we maintain a single source of truth.

This means architecture changes are tracked through git history, documentation stays synchronized with code changes, and the entire team can access it directly in their development environment.

For team members who temporarily lack access to the code repository, we'll generate PDF files from our markdown documentation using **pandoc**. This allows us to share documentation via email until proper repository access is granted ¹⁷.

For more information about pandoc, visit: pandoc.org. The Markdown Preview Enhanced extension also offers PDF generation capabilities from markdown files, which we'll explore in detail later.

3.3.4 Checkpoint: 004-setting-up-a-clean-workspace

At this point, we should have the state of:

github.com/nxpatterns/elvis/tree/004-setting-up-a-clean-workspace

¹⁷ In fact, the documentation you're reading right now was rendered from markdown to PDF using **pandoc**, demonstrating this exact workflow in practice.

4 NestJS Setup

4.1 Preparing Nx for NestJS

First, we shall equip our Nx workspace with NestJS superpowers. To accomplish this, we'll install the necessary packages as development dependencies (hence the `-D` flag), starting with:

```
1 npm i -D @nx/nest
```

Additional packages will follow (as we continue our NestJS journey).

4.2 Creating NestJS App Using Nx Generator UI

Now we'll right-click on the apps folder and select "Nx Generate (UI)" from the context menu.

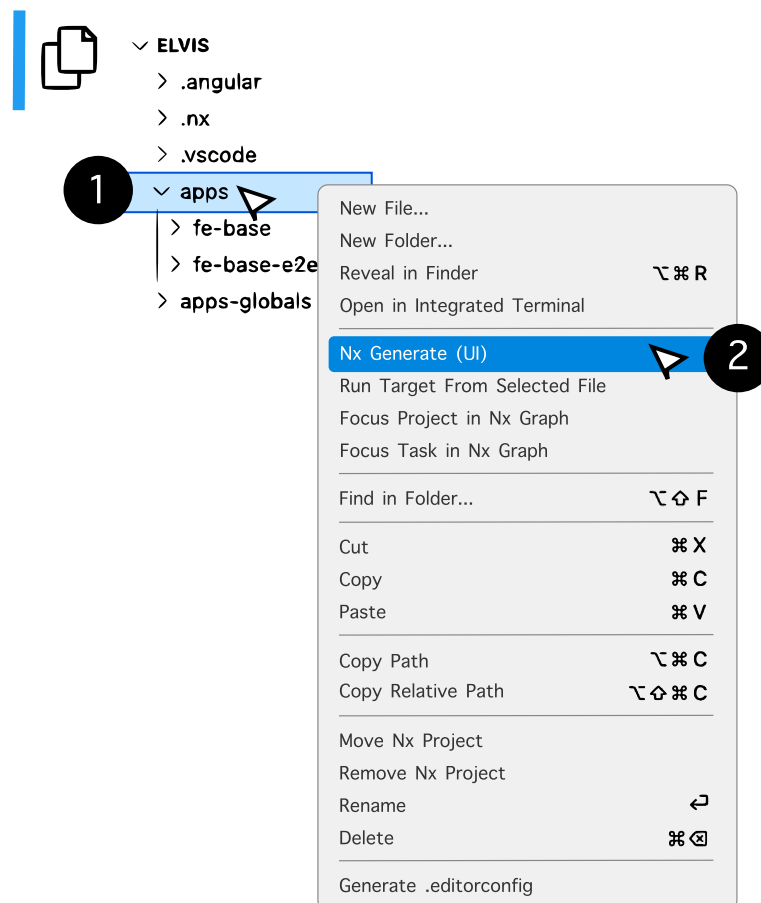


Figure 17: Nx Generator UI: Right-Click on “apps” Folder

This will open the Nx Generator Interface in the editor's right pane where we'll carefully configure our entries, as we already have an Angular frontend app - they'll be communicating with each other soon. First, we need to select which generator to use. Among the many available options, simply type `nx/nest` and select the first entry (note: the order might vary in different versions): “@nx/nest - application Create a NestJS application.”




Figure 18: Nx Generator UI: Selecting “@nx/nest” Application

Subsequently, our Nx Generator will present these configuration options:

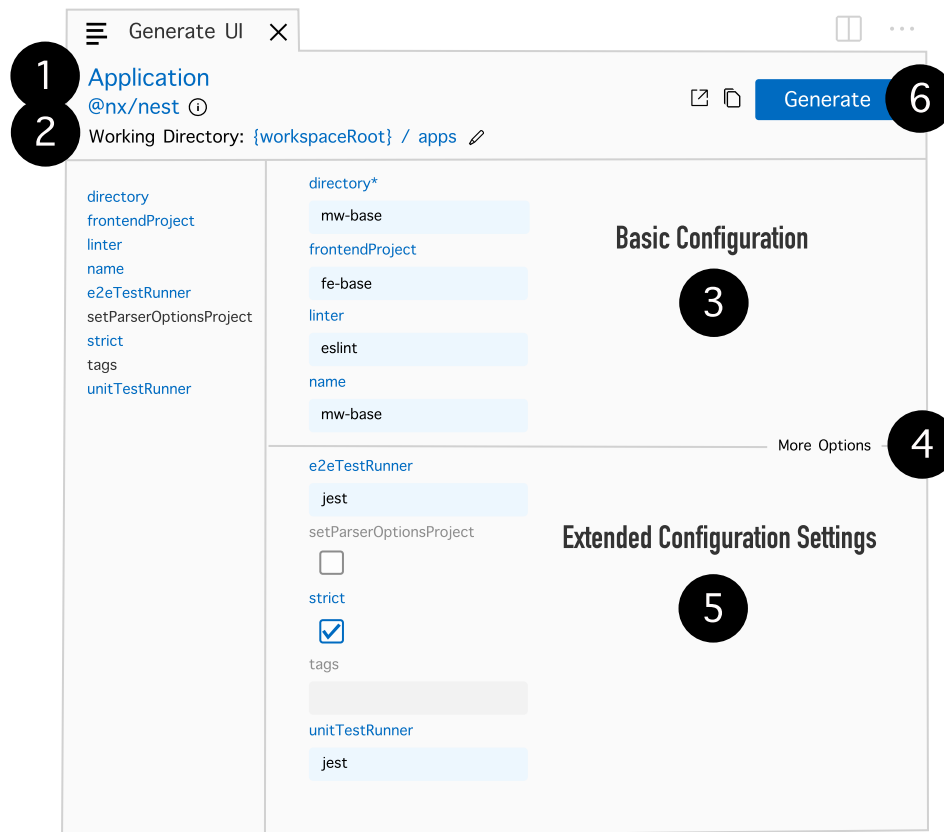


Figure 19: Nx Generator UI: NestJS App Configuration

Let’s examine which configurations we’ll apply and why:

4.2.1 Application Type (@nx/nest)

First, we verify that we’re using the **@nx/nest** generator. There are numerous generators available, and NestJS itself has one. The configuration options often differ, sometimes significantly. By confirming the correct generator first, the subsequent values should be accurate (while newer versions might have different values, they should look similar in principle, as the underlying concepts don’t change frequently).

4.2.2 Working Directory

In our current approach, we initiated app generation through the VSCode UI. We had installed the Nx Console Extension, which integrated its own entries into the context menu.

In the Explorer view (Explorer icon selected in the Activity Bar), we navigated to our desired apps directory, right-clicked to open the context menu, and selected “Nx Generate UI”.

This intention must be reflected in the generator configuration. Therefore, the selected Working Directory must match `{workspaceRoot}/apps`. We verify this here. If it doesn’t match, we can correct it using the adjacent pencil icon.

4.2.3 Basic Configuration

You’ll notice that point 3 refers to an entire section and we have four fields: **directory**, **frontendProject**, **linter**, and **name**. This represents the basic configuration, as further down you’ll see a clickable “Show all Options” button that reveals “Extended Configuration Settings”.

4.2.3.1 Directory

Here, I select the app name as the directory name - I’ll explain the rationale behind this choice in the Name section below.

4.2.3.2 Frontend Project

We’ll select ‘fe-base’ here, as it’s our existing Angular frontend application. Since this guide is written for Angular experts venturing into NestJS territory, connecting our frontend makes perfect sense.

This connection is essential as it configures the necessary proxy settings for our development environment. Later, when we start testing the communication between our frontend and middleware, this configuration will ensure proper routing of requests. If you’re not planning to work with a frontend immediately, you could leave this field empty - though it’s easier to set it up now rather than later.

4.2.3.3 Linter

For our linting configuration, we’re selecting ESLint - the de facto standard for TypeScript projects.

While both ESLint and other linting tools have their merits, ESLint’s extensive ecosystem and robust TypeScript support make it the natural choice. As Angular experts, you’ll appreciate that this aligns with Angular’s own tooling choices, ensuring consistency across your full-stack application.

4.2.3.4 Name

Project names are typically set by product management. I prefer adding only a prefix like fe-, mw-, be-, etc. Here, we’ve chosen “mw-base” as I’m planning to write an “Elvis” book that will explore data modelling practices as a foundation for machine learning in detail. We’ll use this “mw-base” app as a foundation and later generate other Elvis-compatible apps using either an Nx plugin or a VSCode extension. If you’re creating your own NestJS app, you’re free to choose any name you prefer.

4.2.4 Extended Configuration Settings

4.2.4.1 E2E Test Runner

Jest is our choice for end-to-end testing, providing a familiar testing environment that integrates smoothly with our NestJS backend. I suspect our Angular experts won't mind - after all, Jest's brilliant test isolation and snapshot capabilities make a rather compelling case, don't they?

While we use Playwright for frontend E2E testing (simulating user interactions in the browser), backend E2E testing has a different focus: it tests the complete request-response cycle through our API endpoints, including middleware processing, database operations, and external service integrations. Jest excels at this type of testing, allowing us to simulate HTTP requests and verify our entire backend processing chain works correctly.

50 Shades of E2E Testing

Don't be surprised that our familiar frontend term "E2E" takes on a different context here. There's also another type of E2E that tests FE and BE simultaneously. To thoroughly confuse you:

- Frontend E2E (with or without backend. Without backend -> mocks, stubs, spies, service workers, fake REST APIs, mock servers, test doubles, in-memory databases)
- Backend E2E (with or without frontend)
- Complete Frontend E2E (including SSR) and all connected backends (production-like E2E tests across the entire infrastructure landscape)
- The previous E2E plus distributed networks (locations, synchronisation, etc.) including Kubernetes scaling
- And there are about 10 more types of E2E tests.

Stay tuned for the Elvis book :)

4.2.4.2 Parser Options

We leave this **unchecked** as we're handling parser options in our base TypeScript configuration.

When building a monorepo with multiple applications and libraries, TypeScript's parser can sometimes become confused about project references and file paths. The `setParserOptionsProject` flag, if checked, would create a separate parser configuration for this specific application.

However, we're opting not to use it since we've already established a robust base TypeScript configuration at the workspace root level. This configuration handles all our parser settings, including path aliases and project references, in a centralised manner.

🕒 To Parse, or Not to Parse - That's Configurable

Using the workspace-level configuration is not only cleaner but also ensures consistency across all our applications and prevents potential conflicts that could arise from having multiple parser configurations.

Our approach aligns with the monorepo best practices where we maintain core configurations at the root level, allowing all projects to inherit these settings uniformly. This becomes particularly important when we start generating additional Elvis-compatible applications later.

However, I can assure you that in every monorepo, there will inevitably be that one special application requiring custom parsing - courtesy of some rather “brilliant developers”. But that’s a story for another chapter, isn’t it?

4.2.4.3 Strict Mode

We want to use TypeScript’s strict mode for better type safety and error detection. While Angular developers are quite familiar with this concept, let’s clarify its importance: Strict mode enforces rigorous type checking rules, catching potential issues during development rather than at runtime. It activates several crucial TypeScript flags such as `strictNullChecks`, `noImplicitAny`, and `strictFunctionTypes`.

For those newer to Angular: imagine strict mode as your extremely detail-oriented colleague who reviews your code in real-time, catching type mismatches, preventing null reference errors, and ensuring you’ve thought through all possible scenarios. While it might seem a bit demanding at first (your IDE will certainly let you know!), it’s an invaluable guardian for maintaining code quality in large-scale applications.

This becomes particularly crucial in our middleware layer where data validation and type consistency are paramount for reliable API communication.

Later, when developing Elvis Enterprise Applications, we’ll discover that certain errors are automatically corrected, and faulty code cannot be committed or pushed at all. This check is one of the options that will be crucial for future auto-correction mechanisms. Elvis is designed to improve and test various human errors through automatic corrections (and I don’t even trust myself in this regard).

4.2.4.4 Tags

Left empty for now - this seemingly simple field actually plays a significant role in larger monorepo architectures. Tags help us categorise and organise our applications and libraries, particularly useful when your workspace grows to dozens or hundreds of projects.

For example, you might later tag applications with their domain (‘customer-facing’, ‘admin’, ‘reporting’), their type (‘app’, ‘lib’, ‘utility’), or their status (‘stable’, ‘experimental’, ‘deprecated’). These tags then become powerful tools for:

- Selective building and testing
- Access control and dependency management
- Automated deployment workflows

— Documentation generation

For our current middleware application, we'll keep it untagged until we establish our tagging strategy later. As your Elvis application ecosystem grows, you'll appreciate how tags can help maintain order in complexity.

Think of tags as library catalogue cards in a vast digital library - they help you find and manage related pieces efficiently. For now, we're just placing our first book on the shelf, but we're keeping the option open for future organisation.

4.2.4.5 Unit Test Runner

We've chosen Jest for unit testing, just as we did for E2E tests - and yes, dear Angular developers, Jest is indeed that versatile! In the modern Angular ecosystem, we're already familiar with robust testing approaches, and you'll find Jest's capabilities in the NestJS context both powerful and elegant.

In the NestJS context, unit tests focus on individual components like controllers, services, and pipes in isolation. Think of testing a service's business logic without actually hitting the database, or verifying a controller's response transformation without making real HTTP calls. Jest's exceptional mocking capabilities shine here - whether you're simulating database responses, mocking external services, or testing error scenarios.

E2E tests, as we discussed earlier, examine the complete request flow through your API endpoints. The beauty of using Jest for both testing types lies in its consistent syntax and shared test utilities. You'll write your unit tests in `*.spec.ts` files right next to your implementation files, while E2E tests live in a dedicated `e2e` directory, but both leverage the same powerful testing framework.

For instance, you might unit test a `UserService`'s password hashing method in isolation:

```
1 expect(await userService.hashPassword('secret')).toHaveLength(60);
```

While your E2E test would verify the entire registration flow (Caution - this is pseudo-code; we'll implement this more professionally later):

```
1 const response = await request(app.getHttpServer())
2   .post('/auth/register')
3   .send({ username: 'elvis', password: 'elvis-rocks' });
4 expect(response.status).toBe(201);
```

4.2.5 Completing Configuration

After making these entries, we shall click the “Generate” button to create our first NestJS application in this Nx workspace.

4.2.6 Checkpoint: 005-create-a-nestjs-app

At this point, we should have the state of:

github.com/nxpatterns/elvis/tree/005-create-a-nestjs-app

4.3 NestJS Setup Review

Once you’ve set up an Nx workspace, you quickly discover that everything else runs very smoothly and efficiently. (Don’t worry, I’m not constantly praising the Nx team - BTW, they know me more as someone who asks uncomfortable questions, he he...)

If you now click on the “Nx Console icon” in the “Activity Bar”, you should see our projects (apps): “fe-base” and “mw-base”, each with their respective “E2E” test projects:

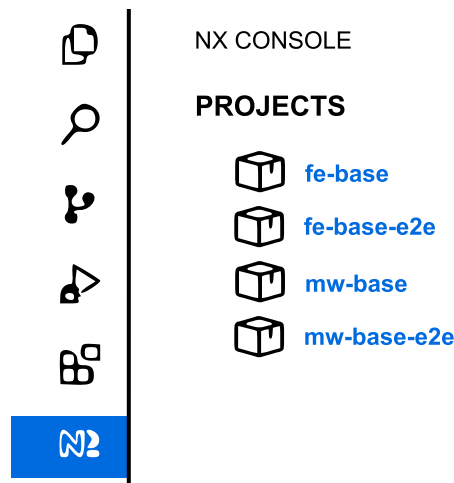


Figure 20: Nx Console: Angular and NestJS Projects

4.3.1 Serving the NestJS Application

Now, just as we did with our Angular app, we'll start our NestJS app and examine its initial output. Click on the **Nx Console** icon in the **Activity Bar** to open the Nx Console.

Then go to **Projects** → **Project Name** (in our case **mw-base**) → **serve** → **development** and click “Execute task” (Play Button).

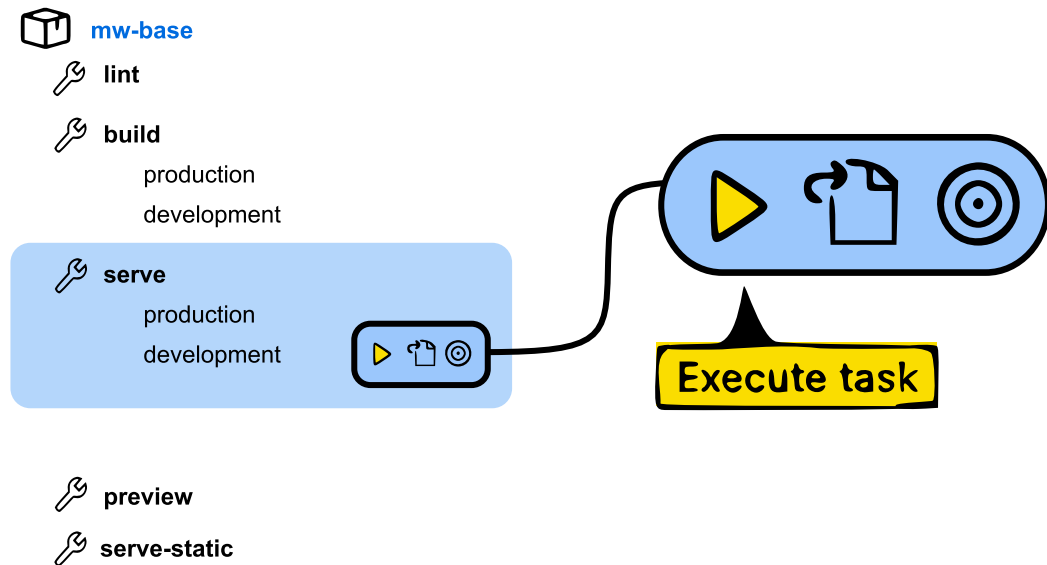


Figure 21: Nx Console Extension: Serve NestJS App

This will initiate a terminal session and start the NestJS development server. Let's examine the terminal output (*excerpt*):

```

1  * Executing task: npx nx run mw-base:serve --configuration=development
2  NX   Running target serve for project mw-base and 1 task it depends on:
3
4  > nx run mw-base:build:development
5  > webpack-cli build node-env=development
6
7  NX Successfully ran target build for project mw-base (37ms)
8
9  Debugger listening on ws://localhost:9229/1088f5f9-e82a-4f97-ae10-a1b51585fd5c
10 For help, see: https://nodejs.org/en/docs/inspector
11
12 [Nest] 17132 - 01/16/2025, 12:05:54 PM LOG [NestFactory] Starting Nest application...
13 [Nest] 17132 - 01/16/2025, 12:05:54 PM LOG [InstanceLoader] AppModule dependencies initialized +5ms
14 [Nest] 17132 - 01/16/2025, 12:05:54 PM LOG [RoutesResolver] AppController {/api}: +2ms
15 [Nest] 17132 - 01/16/2025, 12:05:54 PM LOG [RouterExplorer] Mapped {/api, GET} route +1ms
16 [Nest] 17132 - 01/16/2025, 12:05:54 PM LOG [NestApplication] Nest application successfully started +1ms
17 [Nest] 17132 - 01/16/2025, 12:05:54 PM LOG 🚀 Application is running on: http://localhost:3000/api

```

🐾 The Tale of Two Welcomes

While our frontend friends craft beautiful, engaging welcome pages (and rightfully so - user experience matters!), backend services prefer a more subtle approach. Instead of a grand homepage, we'll get a modest "Hello API" message - because even APIs deserve to say hello, just in their own minimalist way.

Think of it as your application's personality: frontend apps are like warm, welcoming hosts at a party, while backend services are like efficient party planners working behind the scenes - both essential, both amazing at what they do!

After all, in our full-stack world, it's this beautiful dance between presentation and functionality that makes modern applications truly shine ✨

Here, we'll initially focus on these two lines:

```
1 Executing task: npx nx run mw-base:serve --configuration=development
2 LOG 🚀 Application is running on: http://localhost:3000/api
```

The first line shows the Nx task being executed, while the second indicates that our NestJS application is successfully running on port 3000 under the `/api` path.

We'll now edit our `package.json` to include a new script. These scripts will be particularly important for cloud deployment later.

📘 Nx Serve Scripts

For those who started reading with this chapter: Please refer to the "Serving the Angular App" chapter for details, where I discussed naming conventions and explained the roles of "npx" and "nx".

We'll consistently maintain these conventions throughout our Nx workspace and later implement automatic mechanisms to verify and, where necessary, automatically correct them. We'll also install additional useful extensions that can trigger certain automated processes with simple key combinations and remember our sessions.

After editing, the scripts section should look like this, where we've added the `mw-base:serve:development` task.

```
1 "scripts": {
2   "fe-base:serve:development": "npx nx run fe-base:serve
   ↪ --configuration=development --host=0.0.0.0",
3   "mw-base:serve:development": "npx nx run mw-base:serve
   ↪ --configuration=development"
4 }
```

The other script `fe-base:serve:development` was added during our Angular setup. Now we can

close all other terminal sessions, open a new terminal and test our script. This should exhibit exactly the same behaviour as if we had started the task via the Nx Console UI.

```
1 npm run mw-base:serve:development
```

Now let's examine the this line and see what this URL delivers:

```
1 LOG 🚀 Application is running on: http://localhost:3000/api
```

Let's open our browser of choice, type “http://localhost:3000/api” into the address bar, press Enter and admire the output:

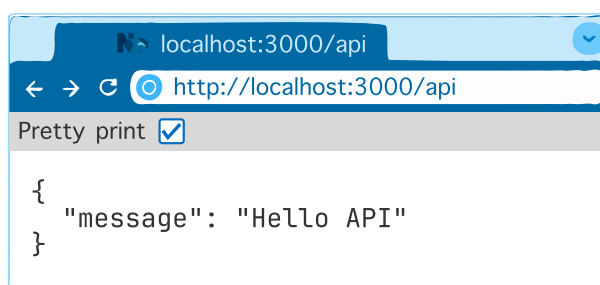


Figure 22: NestJS Welcome Page: “Hello API!”

Now we need to understand where this message comes from and how NestJS works in principle. NestJS follows similar architectural patterns to Angular - with controllers, services, modules, and decorators etc. This architectural similarity is intentional, making the transition particularly smooth for Angular developers moving into backend development.

👁 Before We Move On: Jest Runner Extension

A noteworthy point to add: When we set up NestJS, our recommendations section in `.vscode/extensions.json` gained an additional entry: “firsttris.vscode-jest-runner”. The next time you launch VSCode, it will remind you and ask whether you’d like to install this extension.

4.3.2 Checkpoint: 006-serve-the-nestjs-application

At this point, we should have the state of:

github.com/nxpatterns/elvis/tree/006-serve-the-nestjs-application

5 Architectural Parallels

Before we shift to the NestJS perspective, let's briefly explore the fundamental architectural concepts of Angular and NestJS. Don't worry - I'll keep this dry material as concise as possible while covering all essentials.

Fair Warning: Skip Ahead!

Oh my God! This theory section is quite the sleeping pill (I dozed off while proofreading). Why don't we skip to the hands-on part? It's much more fun! I'll keep this chapter as a perfect example of how not to write a book.

Fear not about skipping this chapter - it won't impact your journey building Enterprise applications with Angular and NestJS in our Elvis Nx monorepo. Only knowledge enthusiasts like myself might struggle to leap forward.

As for historical accuracy - I'm relying on eyewitness accounts rather than first-hand experience. For the technical evolution, best stick to official sources. Though I've checked - they're remarkably quiet on this. Any rumours you hear are my investigative contribution.

A key question: Is this knowledge necessary for Angular and NestJS development? For developers - no. For architects - yes. The reason? Nx generates crucial configurations automatically (like `tsconfig.json` with decorator support) that architects must understand.

Nx simplifies many aspects, but to excel in Enterprise Architecture, familiarity with the following details helps. It's like having pre-made recipes and pre-portioned ingredients - developers can cook perfectly with these, but architects need to know why these specific ingredients and portions were chosen.

5.1 A Natural Partnership: Angular & NestJS

When Angular developers first encounter NestJS, they often experience a moment of recognition. This is no coincidence - NestJS deliberately adopts Angular's architectural principles, creating a seamless transition between frontend and backend development.

Their differences emerge naturally from their distinct responsibilities. Take routing, for instance: NestJS controllers use decorators to define API endpoints, automatically creating a URL structure that mirrors your module hierarchy. A `UserModule` with its `UserController` naturally serves requests under `/users/*`. Angular, focused on client-side navigation, manages routes through a centralized configuration in files like `app.routes.ts`.

Services in both frameworks speak the same language - they're singleton providers, managed through dependency injection, making data flow predictable and testing straightforward.

Let's examine the core architectural concepts - just enough to navigate both frameworks with confidence...

5.2 Decorators & Reflect Metadata API

TypeScript's decorator system relies on two crucial components working behind the scenes. The Reflect Metadata API provides the specifications for storing and retrieving metadata about classes, methods, and their types at runtime. This API is particularly important for frameworks like Angular and NestJS, as it enables their dependency injection systems to work with TypeScript's type information.

However, since the Reflect Metadata API is still a stage 3 proposal, we need the **reflect-metadata polyfill** to make everything work. This is why you'll often see these lines in TypeScript projects:

```
1 import 'reflect-metadata';  
2 // Always import this first, before any other imports
```

Both Angular and NestJS abstract away most of these implementation details, while depending heavily on this infrastructure. However, their approaches to handling these dependencies differ.

For those who compared the last two branches using `git diff`, you might have noticed that creating our NestJS app added a new package `"reflect-metadata": "^0.1.13"` to `package.json`. This explicit dependency is necessary for NestJS projects to function properly.

In Angular projects, this dependency is less visible as it's bundled within `@angular/core` and automatically handled by the framework. You won't find it directly in your `package.json`, but it's an essential part of Angular's dependency injection system.

Angular's Reflect Metadata Evolution

Angular's relationship with Reflect Metadata has evolved significantly, though it's important to note that the core Reflect Metadata API remains fundamental to Angular's architecture. What has changed is Angular's approach to the polyfill implementation.

In production builds, the Ahead-of-Time (AOT) compiler transforms decorators and metadata into highly optimized factory code, eliminating the need for runtime reflection. During development, Just-in-Time (JIT) compilation still utilizes metadata reflection, but Angular handles this internally through `@angular/core`.

The critical distinction is that while Angular still uses metadata reflection principles, it no longer requires developers to explicitly include the `reflect-metadata` polyfill. This architectural decision has notably improved bundle sizes and application startup performance.

This setup enables powerful features like Angular's dependency injection and NestJS's parameter decorators, all while keeping the actual metadata handling invisible to application code. As the JavaScript ecosystem evolves, this foundation might change, but for now, it's a stable and well-tested solution that powers many enterprise applications.

5.2.1 The Decorator Decision: A Pivotal Moment in Angular's History

The transition from Angular 1 to Angular 2 marked a significant architectural evolution in the framework's history. The introduction of decorators represented more than just a technical enhancement; it symbolized a fundamental shift in how Angular approached component architecture and dependency injection.

During the modernization phase, the Angular team faced the complex challenge of reimagining their framework for contemporary web development. The adoption of TypeScript brought static typing and improved tooling, while decorators emerged as an elegant solution for metadata management and component configuration. This decision aligned with emerging JavaScript standards and offered developers a more intuitive way to define application structures.

The path to implementing decorators involved significant collaboration between major technology companies. While the public narrative highlighted the technical benefits and architectural improvements, the behind-the-scenes discussions were more nuanced. The TypeScript team's initial hesitation to implement experimental decorators, given their commitment to JavaScript compatibility, led to constructive dialogues about the future of web development tooling.

Eventually, the implementation of experimental decorators in TypeScript proved to be a pivotal moment, enabling Angular's modern architecture while maintaining harmony in the developer ecosystem. This period demonstrated how technical decisions often involve careful diplomacy and compromise among major players in the software industry.

5.2.2 From Boilerplate to Elegant Architecture: Angular's Decorator Journey

For modern Angular experts, features like clear component structure, type-safe input/output bindings, declarative metadata, and simplified template integration are fundamental aspects of the framework. These capabilities, which we now take for granted, are direct benefits of Angular's modern architecture and particularly its decorator system.

However, there's one historical aspect worth remembering: "manual registration." Consider how straightforward dependency injection is in today's Angular - we simply write

```
1 constructor(private http: HttpClient) {}
```

Now imagine if, for each dependency, you had to explicitly register it:

```
1 app.register('DataService', ['HttpClient', function(http) {  
2   return new DataService(http);  
3 }]);
```

This cumbersome approach would significantly impact development productivity. During Angular's modernization, the combination of TypeScript and decorators moved these mechanical processes under the hood, freeing developers from such implementation details.

When developers use decorators in Angular, they do much more than just labeling a class (like using `@Component` to mark a class as a component). These decorators establish a contract -

they define what properties and behaviors this decorated class must have and how they should be implemented. While a component doesn't necessarily need dependencies in its constructor, when it does have them, they must follow specific patterns (such as `private http: HttpClient`).

Interestingly, when we read Angular documentation or books, they rarely delve into the intricate details of how a component decorator is implemented internally. Instead, they focus on teaching us how to use it effectively. Behind the scenes, decorators perform their work with remarkable elegance, quietly handling complex operations without drawing attention to themselves.

From Framework Patterns to Enterprise AI Architecture

The patterns we explore transcend Angular and NestJS, forming the cornerstone of modern enterprise software development. In Elvis, our meta-grammar will serve dual purposes: establishing compatibility rules for applications (enabling configuration of market-specific software suites) and creating an AI foundation for analyzing and generating enterprise applications based on proven patterns.

Understanding these framework principles goes beyond mastery of individual tools - it's about building the foundation for scalable, AI-ready enterprise architectures. This knowledge is essential for designing meta-grammars that empower AI to comprehend, analyze, and generate sophisticated enterprise applications.

Elvis, our Nx workspace, will implement further battle-tested patterns including fractal geometry for infinite scalability and proxy patterns for true enterprise reusability - and much more. Stay tuned for the Elvis book.

5.2.3 Angular's Strategic Evolution

Angular's relationship with decorators mirrors the broader transitions in the TypeScript ecosystem. While leveraging TypeScript's stable decorator system for core functionality, Angular is thoughtfully exploring new directions. This reflects an interesting divergence in the JavaScript world: TypeScript has declared its decorator implementation production-ready, while the official ECMAScript decorator specification remains in Stage 3 of standardisation.

Angular 19 exemplifies this evolution with its signal-based APIs. We observe a gradual shift where signal inputs offer alternatives to `@Input()` decorators, and signal-based queries provide new approaches to what `@ViewChild`, `@ViewChildren`, `@ContentChild`, and `@ContentChildren` traditionally handled. For teams exploring these changes, the framework provides migration tools via `ng generate @angular/core:signal-queries-migration`.

The introduction of the `strictStandalone` compiler flag marks another step in Angular's journey, suggesting alternatives to traditional module-based architecture. Unlike the significant shift from AngularJS to Angular 2, this evolution maintains backward compatibility while opening new possibilities. Angular's parallel development strategy - maintaining established decorator patterns while exploring signal-based approaches - allows teams to adopt new patterns at their own pace with full support for existing applications.

While decorators represent a mature, well-defined pattern with clear implementation rules enforced at build time, the emerging landscape of signals, zoneless architecture, and the new Resource API introduces both opportunities and challenges.

The combination of these features with existing reactive paradigms like RxJS creates a complex decision space for architects. Even with just three variables in play, teams face 2^3 possible implementation patterns, each with its own performance characteristics and trade-offs.

Establishing best practices for optimal runtime performance and scalability in this new ecosystem requires careful evaluation and benchmarking - a critical task that warrants its own detailed investigation beyond the scope of this documentation. Rest assured, I'm committed to exploring this thoroughly, as benchmarking is one of my favorite aspects of software engineering.

5.3 NestJS: The Rise of an Enterprise Node.js Framework

While **Express.js** had been faithfully serving **Node.js** applications since 2010, the enterprise world yearned for something more structured. When Angular 2 emerged in late 2015, brandishing **TypeScript** and sophisticated architectural patterns, it sparked an idea.

Enter NestJS in 2017, cheekily borrowing Angular's best bits - decorators, dependency injection, and modular architecture. Like a well-mannered guest, it didn't just copy; it adapted these patterns brilliantly for backend needs. Initially embracing Express.js as its foundation, NestJS later welcomed Fastify to its party, offering developers a choice of underlying engines.

What started as Angular's backend counterpart quickly evolved into a framework-agnostic powerhouse. Today, it plays nicely with GraphQL, WebSockets, gRPC, and various microservices patterns, rather like a Swiss Army knife wearing a tailored suit.

This TypeScript-native framework brought enterprise-grade architecture to Node.js without the enterprise-grade headaches - a feat that made many Java architects raise their eyebrows, and some even their spirits.

5.3.1 NestJS' Strategic Evolution

NestJS began its journey with a clear mission: bringing Angular's elegant patterns to server-side development. Its initial release focused on core fundamentals - **modules**, **controllers**, and **providers** - essentially the building blocks that made Angular shine.

Each major version brought calculated expansions. Version 4 introduced **custom decorators** and improved dependency injection. Version 5 added **CLI tools**, making project scaffolding as smooth as buttered toast. Version 6 brought **Fastify support**, offering developers a speedier alternative to Express.

The framework's appetite grew with its success. Version 7 embraced **microservices patterns**, while Version 8 refined its **GraphQL capabilities**. Version 9 introduced **improved TypeScript features and enhanced performance**, making it nimbler than a caffeinated squirrel.

By Version 10, NestJS had matured into a complete ecosystem, sporting features from automated **Swagger documentation** to sophisticated **caching strategies**. Version 11, freshly released in January 2025 (github.com/nestjs/nest/releases), continues this tradition of excellence, promising further refinements and innovations, with comprehensive migration guidance (docs.nestjs.com/migration-guide) forthcoming in the official documentation.

The most fundamental change appears to be the discontinuation of Node v16 and v18 support (\geq v20 is now required). Additionally, NestJS v11 has adopted Express v5 and Fastify 5 as

its default versions, introducing several breaking changes. For detailed migration guidance, please refer to:

- Express v5: expressjs.com/en/guide/migrating-5.html
- Fastify v5: fastify.dev/docs/latest/Guides/Migration-Guide-V5/

6 Upgrade NestJS to v11+ in an Nx Workspace

First we install the latest version of Nx/CLI (see the official documentation here: nx.dev/recipes/installation/update-global-installation)

```
1 npm install -g nx@latest
2 nx --version # global version should be v20.3.2 or higher
```

Let's check if there's a migration configuration available for NestJS v11+ that could handle the necessary changes:

```
1 nx migrate latest
```

The output reveals:

```
1 NX The migrate command has run successfully.
2
3 - package.json has been updated.
4 - There are no migrations to run, so migrations.json has not been created.
```

Examining the package.json shows that only the @nx packages were updated to their latest versions:

```
1 "@nx/angular": "20.3.0",      -> "20.3.2",
2 "@nx/devkit": "20.3.0",      -> "20.3.2",
3 "@nx/eslint": "20.3.0",      -> "20.3.2",
4 "@nx/eslint-plugin": "20.3.0", -> "20.3.2",
5 "@nx/jest": "20.3.1",        -> "20.3.2",
6 "@nx/js": "20.3.1",          -> "20.3.2",
7 "@nx/nest": "^20.3.1",       -> "20.3.2",
8 "@nx/node": "20.3.1",        -> "20.3.2",
9 "@nx/playwright": "20.3.0",   -> "20.3.2",
10 "@nx/web": "20.3.0",         -> "20.3.2",
11 "@nx/webpack": "20.3.1",     -> "20.3.2",
12 "@nx/workspace": "20.3.0",    -> "20.3.2",
```

The situation we're facing is hardly surprising. With the ink barely dry on the latest **NestJS v11+** release, **Nx** hasn't had the chance to craft a migration configuration yet.

🕒 Automated Migration: The Ideal Scenario

Had a migration configuration been available, we could have completed all current upgrades with just these elegant commands:

```
1 nx migrate latest           # Check for the latest migration
2 npm install                 # Update all dependencies
3 nx migrate --run-migrations # Apply the migration configuration
4 nx update @nx/workspace     # Update the Nx workspace
```

This means we're in for a bit of an adventure - a manual upgrade process that harkens back to the good old days of software development. But don't worry, we'll tackle this step by step. First things first, let's ensure we're working with the latest dependencies:

```
1 npm i
```

Now, it's time to roll up our sleeves and dive into the upgrade process. To make our lives easier, we'll enlist the help of a powerful ally: **npm-check-updates** (affectionately known as **ncu** in command-line circles). This Swiss Army knife of dependency management is about to become our best friend. Let's get it installed globally:

```
1 npm i -g npm-check-updates
```

This nifty tool will be our guide through the jungle of **package.json**, helping us identify and navigate the available updates for all our dependencies. However, a word of caution: while **ncu** is powerful, it's not infallible.

In certain configurations, it might not only be unhelpful but could potentially disrupt existing setups. This tool requires both careful handling and considerable experience to use effectively.

When we run **ncu** without a parameter, it won't make any changes to your system. Instead, it will analyse the installed packages in **package.json** and suggest updates/upgrades:

```
1 ncu
```

🐾 The Good, the Dev, and the Peer Dependencies

Before proceeding with these updates, let's explore the three types of dependencies in Node.js projects - and you'll spot two of them lurking in your package.json:

- **Dependencies** (in package.json under “dependencies”): These are the packages your application simply can't live without in production. Think of them as the bread and butter of your app - `@nestjs/core` and `@nestjs/common` are like the vital organs keeping your NestJS application alive and kicking.
- **Development Dependencies** (under “devDependencies”): These are your application's gym equipment - only needed during development or in your cloud pipeline. Examples include `@nestjs/testing` for keeping your code fit through testing, or TypeScript for giving your JavaScript some proper manners. They won't be invited to the production party.

Now, here's where it gets interesting - **Peer Dependencies** (“peerDependencies”). Since npm v7+, these chaps install themselves automatically, but in the older days, they needed a manual invitation to the project.

In our case, when upgrading from **NestJS v10** to **v11** without “Nx Migration Configuration” and updating packages with `ncu`, we're likely to encounter some “**Peer Dependency Mismatch**” drama.

What are Peer Dependencies? They're the sophisticated cousins in the Node.js family, particularly crucial for library or plugin developers. Think of them as essential guests at your library's party - packages that your library needs but doesn't bring along. Instead, it assumes the host (your project) has already sent them an invitation.

Looking at our output, these packages stand out as potential challenge points (based on known breaking changes):

| | | | | |
|----|---------------------------------------|----------------------|----------------|----------------------|
| 1 | <code>@nestjs/common</code> | <code>^10.0.2</code> | <code>→</code> | <code>^11.0.3</code> |
| 2 | <code>@nestjs/core</code> | <code>^10.0.2</code> | <code>→</code> | <code>^11.0.3</code> |
| 3 | <code>@nestjs/platform-express</code> | <code>^10.0.2</code> | <code>→</code> | <code>^11.0.3</code> |
| 4 | <code>@nestjs/schematics</code> | <code>^10.0.1</code> | <code>→</code> | <code>^11.0.0</code> |
| 5 | <code>@nestjs/testing</code> | <code>^10.0.2</code> | <code>→</code> | <code>^11.0.3</code> |
| 6 | <code>@types/node</code> | <code>18.16.9</code> | <code>→</code> | <code>22.10.7</code> |
| 7 | <code>eslint-config-prettier</code> | <code>^9.0.0</code> | <code>→</code> | <code>^10.0.1</code> |
| 8 | <code>eslint-plugin-playwright</code> | <code>^1.6.2</code> | <code>→</code> | <code>^2.1.0</code> |
| 9 | <code>prettier</code> | <code>^2.6.2</code> | <code>→</code> | <code>^3.4.2</code> |
| 10 | <code>typescript</code> | <code>~5.6.2</code> | <code>→</code> | <code>~5.7.3</code> |
| 11 | <code>webpack-cli</code> | <code>^5.1.4</code> | <code>→</code> | <code>^6.0.1</code> |

6.1 Checkpoint: 007-upgrade-nestjs

At this point, we should have the state of: github.com/nxpatterns/elvis/tree/007-upgrade-nestjs

6.2 The “ncu -u” Upgrade Experiment

With a dash of daring, let’s venture into the unknown by running `ncu` with the `-u` parameter - the ultimate `package.json` adventure:

```
1 ncu -u
```

And voilà - we’re now the proud owners of an updated `package.json`. Buckle up! The `-u` flag automatically updates your `package.json` rather than just suggesting changes. It’s like skydiving without checking your parachute first - thrilling but requires attention!

Since we’re experimenting (though waiting for the official migration configuration would be more sensible - but where’s the fun in that?), let’s create an experimental branch that we can jettison if things go sideways:

```
1 git checkout -b 007A-upgrade-nestjs-experiment-with-ncu
```

Let’s continue our delightful descent into chaos by running “`npm i`” - hold onto your keyboards and embrace the madness! Think of it as a roller coaster ride through dependency hell - eyes closed, screaming optional. 🎢 (output simplified and shortened for readability).

```
1 > npm i
2
3 npm error code ERESOLVE could not resolve
4 npm error While resolving: @elvis/source@0.0.0
5 npm error Found: @angular-devkit/build-angular@19.0.6
6 npm error node_modules/@angular-devkit/build-angular
7 npm error   dev @angular-devkit/build-angular@"~19.1.3"
8 npm error   peer @angular-devkit/build-angular@"
9     ≥ 17.0.0 < 20.0.0" from @nx/angular@20.3.2
10 npm error Conflicting peer dependency: @angular/compiler@19.1.2
11 npm error     2 more (@angular/localize, ng-packagr)
12 npm error Fix the upstream dependency conflict, or retry
13 npm error this command with --force or --legacy-peer-deps
14 npm error to accept an incorrect (and potentially broken) dependency resolution.
15 npm error
16 npm error For a full report see:
17 npm error ~/.npm/_logs/...-eresolve-report.txt
18 npm error A complete log of this run can be found in: ~/.npm/_logs/...
```

Now we face the consequences of our daredevil approach. If you’re seeing “`ERESOLVE could not resolve`” or “`Conflicting peer dependency`” errors - congratulations, you’ve successfully followed our guide on “How to Create Chaos”! 🎯

But fear not, for every foolish move, `npm` provides an equally questionable solution. Enter our magical incantation:

```
1 npm i --legacy-peer-deps
```

Everything installs beautifully, except for those pesky npm security warnings screaming “RED ALERT!” A friendly reminder: This experimental approach is about as production-ready as a chocolate teapot. ☕ Keep this strictly in your lab environment!

```
1 added 42 packages, removed 70 packages, changed 88 packages, and audited 1584 packages in 17s
2
3 232 packages are looking for funding
4   run `npm fund` for details
5
6 3 moderate severity vulnerabilities
7
8 To address all issues (including breaking changes), run:
9   npm audit fix --force
10
11 Run `npm audit` for details.
```

Instead of going nuclear with “`npm audit fix --force`” (which would guarantee a spectacular failure), let’s try the diplomatic approach:

```
1 npm audit fix --legacy-peer-deps # Don't forget the magic flag!
```

Knowing full well it’s likely to be as effective as a chocolate fireguard - but at least the audit report will reveal our sins (and we architects shall document them meticulously) shortened for readability:

```
1 # npm audit report
2 webpack 5.0.0-alpha.0 - 5.93.0 - Severity: moderate
3 webpack's AutoPublicPathRuntimeModule has a DOM Clobbering Gadget
4 that leads to XSS - https://github.com/advisories/GHSA-4vvj-4cpr-p986
5 fix available via `npm audit fix --force`
6 Will install @nx/angular@20.1.4, which is a breaking change
7 node_modules/webpack
8   @nx/module-federation *
9     Depends on vulnerable versions of webpack
10 node_modules/@nx/module-federation
11   @nx/angular ≤0.0.0-pr-29636-e3c31b7 || ≥20.2.0-beta.0
12     Depends on vulnerable versions of @nx/module-federation
13 node_modules/@nx/angular
14 3 moderate severity vulnerabilities
```

We’ve got a classic case of webpack drama with a side of XSS vulnerability with the following key issues:

- Webpack 5.0.0-alpha.0 - 5.93.0 has a moderate severity XSS risk
- The dependency chain: webpack -> @nx/module-federation -> @nx/angular
- Breaking changes ahead with @nx/angular@20.1.4

Spoiler alert: Using `--force` here would be like trying to fix a paper cut with a chainsaw. 🛠️ Let’s hold off on that, particularly with @nx/angular involved.

How would an architect handle such dependency challenges? While a seasoned architect would

typically avoid such entanglements altogether (They would wait for the official migration guide), we find ourselves already deep in dependency complexity. As someone once wisely remarked (or so I believe), when traversing through hell, one must “keep going.”

The situation demands a methodical analysis and resolution. Our assessment reveals that the webpack vulnerability intertwines with `@nx/module-federation`, subsequently impacting our `@nx/angular` configuration. The breaking change warnings indicate potential architectural ramifications that require careful consideration.

A more prudent approach involves sequential updates, beginning with `@nx/angular` at version 20.1.4 using precise version control, followed by modernising `module-federation`, and concluding with webpack’s latest stable release.

```
1 npm install @nx/angular@20.1.4 --save-exact --legacy-peer-deps
2 npm install @nx/module-federation@latest --save-exact --legacy-peer-deps
3 npm install webpack@latest --legacy-peer-deps
```

Running `npm audit fix --legacy-peer-deps` again confirms that our troublesome friend is still present:

```
1 webpack 5.0.0-alpha.0 - 5.93.0
2 Severity: moderate
3 Webpack's AutoPublicPathRuntimeModule has a DOM Clobbering Gadget
4 that leads to XSS - https://github.com/advisories/GHSA-4vvj-4cpr-p986
5 No fix available
6 node_modules/@nx/module-federation/node_modules/webpack
7 @nx/module-federation *
8 Depends on vulnerable versions of webpack
9 node_modules/@nx/module-federation
10 2 moderate severity vulnerabilities
```

Even if you try to fix the root cause in your source files - like line 4 in “apps/mw-base/webpack.config.js” - the error stubbornly persists:

```
1 module.exports = {
2   output: {
3     path: join(__dirname, '../..../dist/apps/mw-base'),
4     publicPath: '/', // <-- This is the culprit
5   },
6   plugins: [
7     new NxAppWebpackPlugin({
8       target: 'node',
9       compiler: 'tsc',
10      main: './src/main.ts',
11      tsConfig: './tsconfig.app.json',
12      assets: ['./src/assets'],
```

```
13     optimization: false,  
14     outputHashing: 'none',  
15     generatePackageJson: true,  
16   }},  
17 ],  
18 };
```

We face a critical decision point regarding our approach. Living with the moderate security issue while setting `publicPath` does mitigate the vulnerability, though `npm audit` will continue to raise alerts. This path requires thorough documentation and explicit risk acceptance. The alternative is waiting for the official Nx migration - a secure approach that aligns with best practices without compromising security.

The Dependency Rollercoaster: A Survival Guide

Think of `npm audit` as an overzealous security guard who keeps flagging you even after you’ve shown your credentials. Documentation here serves as your written permission to pass!

As a Software Architect (if involved from the start), I could largely avoid such issues by strictly adhering to official, secure, validated versions. However, reality rarely aligns with ideals. Sometimes you join mid-project because the previous architect departed, leaving a team desperately seeking solutions.

Sometimes you have “star” developers (with board connections) who lack patience and attempt “coups” during your holiday, trying package updates as we’ve demonstrated here.

And sometimes it’s a Product Manager questioning why we’re not using the latest Angular version. When I explain we’re waiting for official Nx migration configurations, a Business School graduate might challenge the Nx decision itself, noting it in bold red letters in board reports.

Conclusion: You can do everything by the book and still end up on the pyre. The key is to document every decision, every step, every risk, and every potential consequence. This way, when the fire starts, you can at least say, “I told you so.”

Checking our apps, we find they’re running smoothly (even before an official migration config exists):

```
1 npm run fe-base:serve:development # Angular v19.1.2  
2 npm run mw-base:serve:development # NestJS v11.0.3
```

The only catch is our “`@nx/module-federation`”: “`^20.4.0-beta.1`” package, which would require special security clearance for production use.

Waiting for the official migration configuration remains the safest path. Sometimes a small

donation can motivate the open-source community to expedite matters - a route I gladly take in projects where I control the budget.

In extreme cases (like when an unmaintained library has a security vulnerability), you must roll up your sleeves. Ironically, as soon as you fix the vulnerability, the package suddenly becomes maintained again (yes, I've been there!). In such cases, we had to manually add the correct versions to `package-lock.json` and set up the project using the `npm ci` command.

`npm ci` (clean install) is like `npm install` but stricter - it deletes `node_modules` and installs exact versions from `package-lock.json`, ensuring consistent installations across all environments. Perfect for CI/CD pipelines and when you need absolute version control. 🔒

6.2.1 Beyond Migration: Tracking Critical NestJS Updates

Even though we upgraded our NestJS from version 10 to 11 without an existing Nx migration configuration, we're not finished. What about security patches/fixes that were released during this process? The most reliable source for this information is the Release Notes: github.com/nestjs/nest/releases

If you set up a watcher on "github.com/nestjs/nest" and click on custom, you can check the "Releases" checkbox to stay prepared for such cases. **And indeed**, I received an email today with the following content (excerpt):

```
v11.0.5 (2025-01-23) Bug fixes (@nestjs/core):
#14495 fix(core): global module middleware should be executed first.
```

This means we need to apply this fix in our Nx Mono Repo as well. Now, let's execute these two commands:

```
1 nx migrate @nestjs/core@11.0.5 # The version setter charts our course
2 npm i --legacy-peer-deps       # The implementer makes it reality
```

The first command will update `package.json`, setting the correct version, while the second executes the actual installation. The `--legacy-peer-deps` flag is necessary because we manually updated everything using `ncu`. Once a proper Nx Migration Configuration is in place, this flag becomes obsolete.

6.2.2 Checkpoint: 007A-upgrade-nestjs-experiment-with-ncu

If you're brave enough to attempt this `ncu` experiment, your code state should match: github.com/nxpatterns/elvis/tree/007A-upgrade-nestjs-experiment-with-ncu

6.3 From Nest CLI to Nx: Upgrade & Migration Patterns

For the sake of completeness, let's explore another possible migration path. While we won't execute this approach now, it occasionally presents itself as the superior alternative.

Consider our scenario: upgrading NestJS within an Nx repository without an existing migration configuration. Our previous chapter demonstrated `ncu` as an emergency solution—quite

serviceable in many cases, though potentially problematic in others. We established that awaiting official upgrade migration configurations represents the ideal strategy.

Yet what if we find ourselves unwilling to wait, yet hesitant to employ `ncu`? Enter an elegant alternative, predicated on the Nest CLI's native capabilities. One rather expects native CLIs to handle their latest versions properly, and indeed, they typically do.

At the initial project creation stage, we could install the latest Nest CLI version:

```
1 npm i -g @nestjs/cli@latest
```

This enables us to create a pristine NestJS application outside our Nx ecosystem using the command `nest new example-service`. (Naturally, we should select a more descriptive name reflecting the application's purpose—`nest-v11` serves merely as a demonstration placeholder and wouldn't be suitable for production nomenclature).

We'll use exact the same npm setup to maintain consistency with our current Nx configuration, though superior alternatives exist—perhaps worthy of mention later in this document.

After entering this new project directory, executing `nx init` initiates a fascinating dialogue. Nx inquires about cacheable NPM scripts, their additional outputs (such as those residing in `dist` directories), linter preferences, and various other configuration details. Once Nx concludes with its cheerful “NX 🎉 Done!”, we verify our setup with:

```
1 npx nx run-many -t build
```

A successful build manifests thusly:

```
1 } npx nx run-many -t build
2  ✓  nx run nest-v11:build (2s)
3  _____
4  NX   Successfully ran target build for project nest-v11 (2s)
```

Now comes the truly engaging part. Having established proper Nx configuration settings, we can import this project into our Nx workspace.

We typically create our own importers to automate these procedures. Because, while Nx provides its native `nx import` functionality (documented at nx.dev/recipes/adopting-nx/import-project) that can automate many scenarios, our Elvis ecosystem requires additional sophistication.

🕒 The Future of Integration: AI-Ready Import Strategies

Later, in the Elvis book, I'd like to explore crafting custom Nx plugins and VSCode extensions that elegantly handle external application imports (requiring only TypeScript as a foundation). During the import process, we will equip these applications with a sophisticated meta-language—a rather ingenious approach that serves dual purposes:

- First, it orchestrates harmonious configuration across our imported applications, creating a unified development symphony, if you will.
- Second, it establishes a structured linguistic foundation for AI comprehension, paving the way for future automation of these intricate processes.

This meta-language becomes our architectural Rosetta Stone, elegantly bridging human design intent with AI-driven automation possibilities.

For manual import operations, you can simply move the NestJS application into the `apps/` folder and adjust Nx configuration accordingly. We'll explore this topic thoroughly in our forthcoming "Nx Project Configurations" documentation. The official documentation provides exemplary foundational guidance: nx.dev/reference/project-configuration

🕒 Package Updates: The Security Blind Spot

When might this approach prove particularly advantageous? Consider scenarios where `ncu` upgrades an existing package, yet the new application version has deprecated it—perhaps due to security vulnerabilities. Since `ncu` merely analyses and adjusts versions rather than removing packages, we might encounter persistent `npm audit` issues, even with unused packages lingering in `package.json` due to oversight or suboptimal workspace configuration.

Modern build configurations typically employ tree-shaking to eliminate unused elements. However, since our cloud pipeline examines `package.json` and/or `package-lock.json` as its foundation, security teams might raise legitimate concerns.

6.4 Re-Check Versions & Patches

Before diving into NestJS Basic Concepts for Angular developers, let's conduct a thorough workspace audit using `ncu`. Not just for NestJS - our entire dependency landscape deserves scrutiny.

Semver's elegant contract permits - nay, demands - we apply patches and fixes. It's our infrastructure's immune system at work. Our `ncu` scan reveals a satisfying harvest: predominantly fixes and one backwards-compatible minor update.

```
@angular-devkit/build-angular    ~19.1.3 → ~19.1 .4
@angular-devkit/core             ~19.1.3 → ~19.1 .4
@angular-devkit/schematics       ~19.1.3 → ~19.1 .4
@angular/animations              ~19.1.2 → ~19.1 .3
@angular/cli                     ~19.1.3 → ~19.1 .4
@angular/common                  ~19.1.2 → ~19.1 .3
```


| | | | | |
|-----------------------------------|---------|---|---------|----------------|
| @angular/compiler | ~19.1.2 | → | ~19.1.3 | |
| @angular/compiler-cli | ~19.1.2 | → | ~19.1.3 | |
| @angular/core | ~19.1.2 | → | ~19.1.3 | |
| @angular/forms | ~19.1.2 | → | ~19.1.3 | |
| @angular/language-service | ~19.1.2 | → | ~19.1.3 | |
| @angular/platform-browser | ~19.1.2 | → | ~19.1.3 | |
| @angular/platform-browser-dynamic | ~19.1.2 | → | ~19.1.3 | |
| @angular/router | ~19.1.2 | → | ~19.1.3 | |
| @nestjs/common | ^11.0.3 | → | ^11.0.5 | |
| @nestjs/platform-express | ^11.0.3 | → | ^11.0.5 | |
| @nestjs/testing | ^11.0.3 | → | ^11.0.5 | |
| @nx/angular | 20.1.4 | → | 20.3.2 | # Minor Update |
| @schematics/angular | ~19.1.3 | → | ~19.1.4 | |
| @swc/core | ~1.10.8 | → | ~1.10.9 | |
| @types/node | 22.10.7 | → | 22.10.9 | |

After meticulous review, we're cleared to execute anew:

```
1 ncu -u
2 npm i --legacy-peer-deps
```

Release Vigilance: A DevOps Overture

Still puzzled by our dependency dance with `--legacy-peer-deps` flag? We've waltzed through this particular peer dependency drama in exquisite detail at the chapter's overture. For the fullest appreciation of our technical choreography, might I suggest a graceful pivot to the chapter's beginning?

Think of it as joining a particularly sophisticated ballet mid-performance - while one might grasp the immediate movements, the full artistic context proves rather illuminating. Our earlier exposition provides the complete libretto, as it were, of this dependency management opus.

A touch theatrical perhaps, but then, enterprise architecture does occasionally demand its dramatic moments, wouldn't you say? The chapter's genesis awaits your scholarly attention.

For our newly-joined readers - a brief but crucial staging note: That peculiar `--legacy-peer-deps` flag? It's our temporary choreography workaround, necessitated by our manual `ncu` updates. Think of it as a graceful compromise until our Nx Migration Configuration matures into its proper form.

6.5 Production Readiness: Battle-Testing Our Upgrade

6.5.1 Current State Analysis

The current state is still unsatisfactory. While we've installed the latest versions, the security vulnerabilities remain unaddressed. From the second method covered in the subchapter "From Nest CLI to Nx: Upgrade & Migration Patterns", we learned that NestJS 11+ no longer uses Webpack. However, since we updated our current Nx Workspace using `ncu`, we still have the old Webpack configuration in place. Now we need to correct all of this from a production-

security perspective.

6.5.2 Official vs. Custom Migrations

As we've mentioned, waiting for an official Nx Migration Configuration is the best practice. However, there are cases where our current approach - which we're demonstrating - might become absolutely necessary (in exceptional cases). This is because we want to face real-world conditions and challenges of Enterprise Architecture Management.

A key insight about enterprise migrations: While “waiting for official tooling” sounds safe on paper, it's like waiting for perfect weather to sail - sometimes you need to navigate through the storm! Our approach might seem unconventional, but it reflects the real-world scenarios where business needs don't always align with framework release schedules.

6.5.3 Our Enterprise Migration Manifest

Picture this: We're preparing for a high-stakes theatrical performance, and these are our non-negotiable production requirements. Like any good drama, we need to ensure our show doesn't just look good in rehearsal, but stands up to the bright lights of production!

- Our Custom Upgrade must be as secure as a bank vault on steroids
 - No security vulnerabilities lurking in the shadows
 - Think “Fort Knox” rather than “cardboard castle”
- Our Custom Upgrade must embrace the new configuration standards
 - Out with the old Webpack ways, in with the new NestJS glory
 - No clinging to deprecated standards like a developer with legacy code
- Our Custom Upgrade must play nicely with future official upgrades
 - We're setting the stage for Nx's grand entrance
 - No special snowflake configurations that'll break when the official upgrade arrives
 - Zero “we did it our way” complications that'll haunt us later

Think of it as preparing a five-star restaurant kitchen: We're not just making dinner for tonight; we're setting up systems that'll work seamlessly when Gordon Ramsay drops by for an inspection!

6.5.4 Necessary Steps: The Security-First Approach

First, we need to examine how Nest/CLI has updated the standalone NestJS Application configuration in v11. We'll analyze the current configuration of your Nx workspace and compare it with a standalone NestJS v11 project.

Fortunately, we have identified a single critical update: the `nest-cli.json` configuration that needs to be integrated into our Nx workspace and we have to remove the old configurations.

6.5.4.1 Remove Old Webpack Configuration

Remove `webpack.config.js` from the `mw-base` application:

```
1 rm apps/mw-base/webpack.config.js
```

Remove all webpack-related packages from package.json and install the recommended packages from standalone NestJS v11. I have developed custom Python tools for all these comparison operations and don't worry - you don't need to install Python now and replicate this; you can see the final result here: github.com/nxpatterns/elvis/blob/008-nestjs-basics/package.json. Some of these Python wizardry scripts have found their cozy home here: github.com/nxpatterns/enterprise-software-architecture/tree/main/scripts

6.5.4.2 Integrate nest-cli.json

In enterprise security architecture, maintaining framework compatibility while ensuring robust security measures is critical. The integration of `nest-cli.json` represents a strategic configuration decision that supports both our immediate security requirements and future framework compatibility.

While Nx will likely develop its own configuration standards, our current priority is maintaining secure compatibility with NestJS CLI. This configuration resides at `apps/mw-base/nest-cli.json`:

```
1 {
2   "$schema": "https://json.schemastore.org/nest-cli",
3   "collection": "@nestjs/schematics",
4   "sourceRoot": "src",
5   "compilerOptions": {
6     "deleteOutDir": true
7   }
8 }
```

Security Considerations:

- Schema validation ensures configuration integrity
- Strict source root definition prevents path traversal vulnerabilities
- Controlled output directory management mitigates build-time security risks

This implementation maintains a security-focused bridge between Nx workspace architecture and NestJS framework standards, ensuring a stable foundation for future security patches and framework updates.

Note: This configuration will be reviewed and potentially modified when Nx releases official security-validated migration paths.

6.5.4.3 Update NestJS' project.json

After that, we must modify `apps/mw-base/project.json` to reflect our webpack-free architecture. While I've developed a Python script for this transformation, the comparison and improvement operations are too complex to detail in this documentation.

Instead, let's examine the optimized configuration that ensures compatibility between Nx workspace and NestJS frameworks. You can inspect the generated file here: github.com/nx-patterns/elvis/blob/008-nestjs-basics/apps/mw-base/project.json

Let's focus on the key configuration changes in this refined project.json. The critical transformations are:

- Build executor changed to `@nx/js:tsc` from the previous webpack configuration
- Serve executor updated to `@nx/js:node`
- Simplified test configuration, removing webpack-specific parameters

These changes align with NestJS v11's modern architecture while maintaining the integrity of our Nx workspace configuration. The modifications ensure a clean, webpack-free implementation without compromising our security standards.

6.5.4.4 The Surgical Strike: Security Audit

Now we're left with the final challenge - our pesky vulnerability issue.

The Last Resort Playbook: Emergency Protocols

Let me emphasize this crucial point: You should never find yourself in this position. However, in the real world, companies often rely on third-party libraries that haven't seen updates or maintenance for years. This guide outlines emergency procedures for handling security vulnerabilities in such dependencies.

To be crystal clear: *I don't recommend either custom updates with `ncu` or creating local package copies under normal circumstances.* But sometimes, when you're stuck between a rock and a legacy codebase, there's simply no other way out.

Running

```
1 npm audit fix --legacy-peer-deps
```

again reveals an interesting plot twist:

Based on the output, webpack is still lurking in our project like a stubborn legacy system, specifically in the `@nx/angular` and `@nx/module-federation` packages. It's like finding out your modern microservices architecture still depends on that one COBOL system nobody wants to talk about!

- The `@nx/module-federation` package insists on using vulnerable webpack versions (it's like that one teammate who refuses to update their IDE)
- `@nx/angular` depends on `@nx/module-federation`, creating a delightful dependency chain - think of it as a corporate org chart where everyone reports to the person with the legacy software!

After deeper investigation, we've identified the critical security concern: The security vulnerability stems from a dependency chain where `@nx/angular` depends on `@nx/module-federation`,

which uses an outdated webpack version (5.88.0). This version contains known security vulnerabilities, as identified in our audit warnings. For reference, the current stable webpack version is 5.89.0 (as of January 2025).

To mitigate this security risk, we must implement the following solution:

- Create a local copy of the `@nx/module-federation` package
- Update webpack version in the local package without modifying any code
- Integrate the modified package into our project

This approach minimizes risk by only updating the dependency version without introducing any code changes that would require additional testing. It's a surgical fix that addresses the security vulnerability while maintaining the package's existing functionality and test coverage.

Let's create a local copy of the `@nx/module-federation` package:

```
1 mkdir -p local-modules/@nx/module-federation
2 cp -R node_modules/@nx/module-federation/* local-modules/@nx/module-federation/
```

And then update webpack version in the local package `local-modules/@nx/module-federation/package.json` (excerpt):

```
1 "dependencies": {
2   "webpack": "^5.89.0",
```

Now we have to update workspace `package.json` to reference local package:

```
1 "@nx/module-federation": "file:local-modules/@nx/module-federation",
```

Finally, we remove `node_modules` and `package-lock.json` completely and run `npm install` to apply the changes:

```
1 rm -rf node_modules package-lock.json && npm install --legacy-peer-deps
```

If we then run `npm audit` again, we should see no vulnerabilities:

```
1 > npm audit
2 found 0 vulnerabilities
```

You can see all the changes in this implementation branch: referenced at github.com/nx-patterns/elvis/tree/008-nestjs-basics. I've logged a security issue on Github. By the time this document is complete, a fix may be available.

This temporary mitigation will remain in place until Nx provides an official security update. At

that point, we can safely remove our local modification and transition to the official package.

Current status: **Production-ready** with implemented security patches.

6.6 Checkpoint: 008-nestjs-basics

💡 Cross-Framework Security: Applying NestJS Lessons to Angular

While upgrading NestJS, I identified a similar security vulnerability in **@angular/build** - specifically a moderate severity vulnerability in the **vite** package (versions 6.0.0 through 6.0.8), a dependency of **@angular/build**. I mitigated this using the same approach we recently applied to the NestJS upgrade: creating a local copy, patching the vulnerable version, and referencing it in **package.json** as a temporary solution until an official update becomes available.

At this point, we should have the state of:

github.com/nxpatterns/elvis/tree/008-nestjs-basics

7 NestJS: When Angular Developers Go Full-Stack

7.1 Basic NestJS Concepts for Angular Developers

Welcome to our enchanting overture of NestJS concepts! Don't fear the dramatic flair; the next chapter cuts straight to implementation. Feel free to revisit this section anytime; these theatrical metaphors hide architectural gems in their creative staging.

7.1.1 The NestJS Module: A Familiar Face

NestJS inherits the modular DNA of Angular, though with its own server-side twist. While Angular v19 defaults to “`standalone`” components (which are more like self-contained modules than truly standalone), NestJS sticks with traditional modules.

A Word on Components

Now, here's a plot twist for Angular aficionados: You'll spot the word “**components**” scattered throughout NestJS docs like confetti.

But don't let that trigger your Angular reflexes! In NestJS land, “component” is more like a supporting actor - just a fancy way of saying “piece of the puzzle.”

When you read something like “*Async providers are injected to other components by their tokens*” think “building blocks” rather than Angular components. Everything's a component in this theatrical production - it's just part of the show!

Picture NestJS as Broadway for your backend, where every production starts with a `@Module()` decorator - the playbill that tells Nest exactly what's about to unfold. This isn't just any decorator; it's the master orchestrator that transforms a simple class into a carefully coordinated theatrical universe.

At the heart of every NestJS production is the **root module** - think of it as your opening act. But here's where it gets interesting: Nest uses this to construct something called the “**application graph**.” Imagine a backstage diagram showing how every performer, prop, and scene connects - that's your application graph in action, a masterpiece of architectural choreography!

The Power of Modules: From Stage to Screen

Think of modules as your application's playbills^a - they don't just list the cast and crew (capabilities), they organize them into a coherent, manageable performance that keeps your audience (and developers) coming back for more!

^a A playbill is a program handed out at a theatrical performance, listing the cast, crew, and other essential details about the show.

Now, could you run a show with just one act? Sure, and some small applications do exactly that. But, most real-world applications are more like a full-season theatrical lineup, with

multiple modules each starring in their own specialized performance. Each module becomes a self-contained production, handling its own specific set of capabilities while still being part of the greater show.

7.1.2 Controllers

Speaking of modules, let's meet the star of our backend show: **Controllers**. Don't let the name fool you - while they're not Angular components, these traffic directors are your application's first responders, orchestrating an intricate ballet of incoming requests and outgoing responses.

💡 Gradual Learning: Core Components First

We're keeping **NestJS Middleware** and **NestJS Interceptors** in our back pocket for now - like any good story, you've got to meet the main characters before we introduce the supporting cast!

Every HTTP request that hits your application is like a traveler in need of direction. The routing mechanism is your airport's control tower, determining which controller should guide each request to its final destination. And just like a busy airport handling different types of aircraft, a controller often juggles multiple routes, each designed for a specific kind of request-response tango.

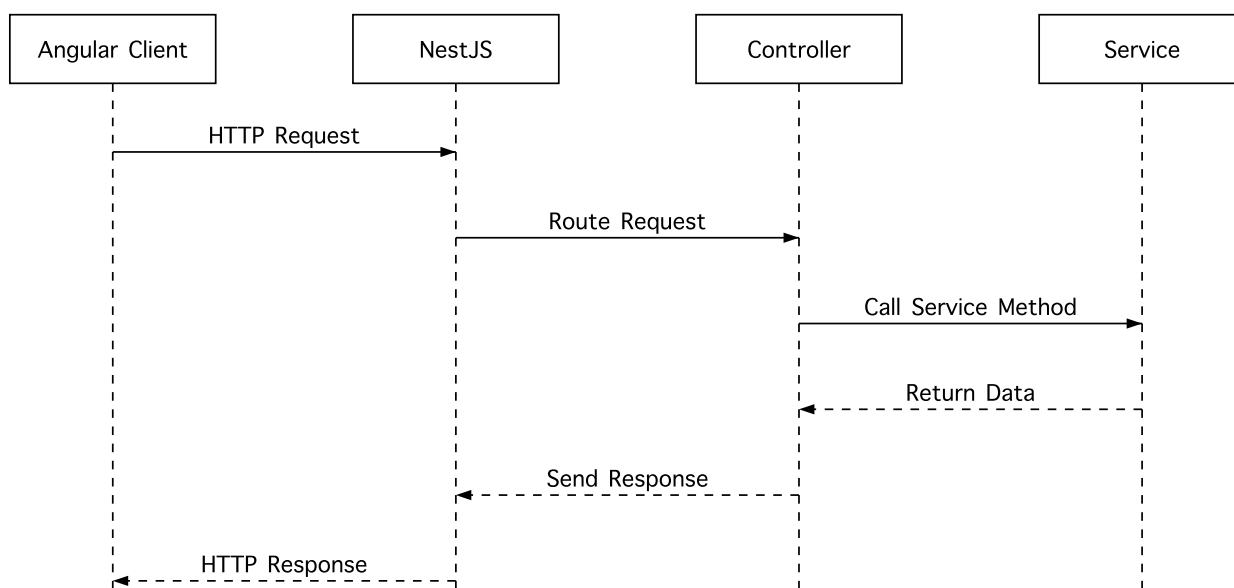


Figure 23: NestJS Core Flow: Basic Request-Response Pipeline

Here's where decorators steal the show - they're the magical staff that transforms ordinary TypeScript classes into **API conductors**.

The `@Controller()` decorator is like pinning a badge on your class, officially deputizing it as a traffic director in your NestJS symphony. These decorators aren't just fancy accessories; they're the secret sauce that helps Nest create its masterful routing map, ensuring every request finds its way home.


```
1 @Controller('cats') // Controller Prefix
2 export class CatsController {
3   @Get('purr') // Route Prefix
4   makeHappySound() { ... }
5 }
```

For our Angular veterans scratching their heads about NestJS’s routing sorcery, fear not! While Angular demands explicit route configurations, NestJS pulls a brilliant architectural sleight-of-hand using decorators and reflection metadata to automatically map routes to controllers. No manual registration required. Think postal service on autopilot:

- Decorators = Building addresses
- Controller methods = Apartment numbers
- NestJS = The all-knowing postal worker



Figure 24: Your request’s journey

It’s like having a GPS that knows every route in town without needing a map. While Angular components might get their spotlight through routes, selectors, or dynamic choreography, NestJS controllers prefer a more direct approach - they’re always on duty, ready to direct traffic with their decorator-powered authority! Your Angular routing anxiety can now officially take a vacation! 🌂

Watch your “NestJS Application Bootstrap Logs” - they’ll reveal all registered routes like a well-organized treasure map:

```
1 # Excerpt from NestJS Application Bootstrap Logs
2 LOG [RoutesResolver] AppController {/api}: +2ms # API Root
3 LOG [RouterExplorer] Mapped {/api, GET} route +1ms
4 LOG 🚀 Application is running on: http://localhost:3000/api
```

Standardizing API Routes: Towards Automated Compliance

Route design isn't about casual `/api/cats/purr` endpoints. We need industrial-strength standards. While e.g. Zalando's REST guidelines^a offer excellent patterns, we're aiming for automated enforcement. (Spoiler: The Elvis Book will reveal our automated compliance solution.)

^a opensource.zalando.com/restful-api-guidelines/ (Last accessed: January 2025)

7.1.3 Providers

Controllers may direct traffic, but **Providers** are the ones making the real magic happen backstage. In NestJS, if something can be injected as a dependency, it's a Provider - think of it as the framework's universal donor system. **Services**, our old friends from Angular, are just one flavor in this dependency cocktail. You've also got **factories** and **helpers** waiting in the wings (more on that stellar cast later).

Providers can be asynchronous

Picture this - your application is like a Broadway show. Sure, the doors are scheduled to open at 8 PM, but you wouldn't start the performance while the orchestra is still tuning up, would you?

That's where async providers shine. They ensure everything's perfectly orchestrated before the curtain rises. Need to establish a database connection before accepting requests? Async providers have got your back.

Providers sync their lifespan with either app-wide or request-bound scopes. During bootstrap, the dependency resolver instantiates each provider like clockwork. App shutdown? All providers get cleaned up - unless they're **request-scoped**. These special providers live and die with their specific request, breaking free from the app's lifecycle. Scopes are defined through decorator parameters, like a VIP pass for your dependencies:

```
1 @Injectable({ scope: Scope.REQUEST })
```

The Scope enum packs three power levels (as of this doc's creation):

- **DEFAULT (0)**: The OG singleton lifestyle
- **TRANSIENT (1)**: Fresh instance every time, like a factory on steroids
- **REQUEST (2)**: New instance per request pipeline - think of it as a request's personal butler

💡 Quick Pipeline Primer

A request pipeline is your HTTP request's journey through middleware city - auth checks, validations, transformations, and business logic, all lined up like a well-oiled assembly line. Each request gets its own express train through this pipeline, keeping things clean and isolated.

We're dissecting the core components like a technical surgeon and afterwards we'll iteratively update our sequence diagram to map the HTTP request's grand adventure with all new details.

7.1.4 Middleware

Remember our "How to Pitch NestJS to Your CTO?" chapter? When infiltrating NestJS into enterprise territory, I deploy the term "middleware" - that magical word that everyone nods at knowingly while mentally painting their own distinct pictures.

🐾 Reality check: Middleware Diplomacy

We're adopting NestJS because we need to transform horrifically suboptimal data (being diplomatically polite here) into something resembling sanity. I sell it as "middleware" - not a "proper" backend (think GoLang, Rust, Java, or those prehistoric backends written in languages lost to time) nor a server-side frontend.

It's the Switzerland of our stack - neutral territory that threatens no one. In this scenario, the entire NestJS application is the middleware diplomat.

In truth, NestJS is a full-fledged backend powerhouse in disguise, capable of replacing virtually any legacy system. But we're not here to start a tech revolution - we're data whisperers on a noble quest. Our mission? Transforming chaotic data streams into semantically rich, well-structured information flows.

Why? Because we're performing CPR on mid-sized corporate dinosaurs in the AI era - and this is our defibrillator. Clean data is the express elevator to AI enlightenment.

Don't be fooled by the term "mid-size" - in Germany, known as "Der Mittelstand", it's the engine powering the entire nation. What makes Germany tick isn't just big corporations - it's these Mittelstand companies, and they're rarely just "mid-sized".

Plot twist: NestJS itself uses "middleware" differently.

Middleware functions are pre-route-handler bouncers with VIP access to request and response objects, plus a mysterious `next()` function (your backstage pass to the next middleware). Think of them as request bodyguards, checking IDs before letting requests meet their final handler destiny.

NestJS offers two elegant paths to middleware mastery: the minimalist function route or the sophisticated class approach with its `@Injectable()` decorator and `NestMiddleware` interface implementation. Now, for the interface stripped to its essence:

```
1 export interface NestMiddleware<TRequest = any, TResponse = any> {  
2   use(req: TRequest, res: TResponse, next: (error?: any) => void): any;  
3 }
```

Generic types **TRequest**/**TResponse** let you type-lock your HTTP interactions (default: wild west **any**). **use()** method is your command center - your chance to inspect, transform, or reject the request. **next()** is your “continue journey” button, optionally packed with errors if things go south. Think of it as a high-stakes relay race where each middleware decides whether to pass the baton (**next()**), drop it (**next(error)**), or keep it forever (no **next()** = request dies here).

Express vs. Fastify

Nest middleware rocks the Express-style party dress by default, but here’s the plot twist: Express and Fastify dance to different beats with their middleware signatures. It’s like twins raised by different framework families - same DNA, different table manners!

Let’s peek at their table manners - where Express brings its time-tested three-course serving style, while Fastify presents a modern tasting menu. Both get your middleware from kitchen to table, but with distinctly different serving ceremonies:

```
1 // Express: The traditional signature  
2 (req, res, next) => void  
3  
4 // Fastify: The modern contender  
5 (request, reply, done) => void
```

The NestJS Pit Stop: Swapping Express for Fastify

Swapping Express for Fastify in NestJS? It’s like upgrading your trusty sedan to a sports car - same destination, different engine! 🚗

The magic happens through Fastify’s **FastifyAdapter**, your VIP pass to speed city. Once bootstrapped, NestJS ditches its Express roots and embraces Fastify’s high-performance ways. But here’s the plot twist: Fastify was the fashionably late arrival to the NestJS party, meaning the official docs are still doing the Express dance in many places.

Word of caution: Those Express-dependent recipes? They’re like trying to fuel your Tesla with diesel - it won’t end well. *Your mission, should you choose to accept it, is to seek out the Fastify-flavored alternatives.* Don’t sweat it though - we’re your pit crew on this performance upgrade! 🛠️

7.1.4.1 Middleware Outside the Box

NestJS middleware architecture tells an interesting story of structured control. The architecture introduces a deliberate ceremony that might initially surprise Angular veterans. Instead of placing middleware in the `@Module()` decorator, NestJS opts for a more sophisticated approach through the `configure()` method, requiring modules to implement the `NestModule` interface.

💡 The Angular-NestJS Mindset Shift

Think of these interfaces as grammar rules in a new programming language. While Angular architects might raise an eyebrow at middleware (it's simply not part of the Angular universe), NestJS demands a specific ceremony.

Don't memorize these interfaces yet - they're like architectural sheet music. We'll perform this middleware symphony later, and that's when these abstract notes will transform into beautiful runtime music. Skip this section freely; we'll encore when implementation time comes.

Only for the sake of completeness, I'm dropping these interfaces into our architectural playbook. They're like the secret handshake that opens the middleware door - a backstage pass to the NestJS universe.

The `NestModule` defines what needs to happen, the `MiddlewareConsumer` determines how to apply it, and the `MiddlewareConfigProxy` fine-tunes where it takes effect. The NestJS middleware architecture reveals itself through these three crucial interfaces; the first beauty unlocks the middleware kingdom:

```
1 export interface NestModule {  
2   configure(consumer: MiddlewareConsumer): any;  
3 }
```

The consumer directs the flow of request processing with surgical precision:

```
1 export interface MiddlewareConsumer {  
2   apply(...middleware: (Type<any> | Function)[]): MiddlewareConfigProxy;  
3 }
```

The proxy is your fine-grained control panel for route orchestration:

```
1 export interface MiddlewareConfigProxy {  
2   exclude(...routes: (string | RouteInfo)[]): MiddlewareConfigProxy;  
3   forRoutes(...routes: (string | Type<any> | RouteInfo)[]): MiddlewareConsumer;  
4 }
```

Now, putting theory into practice, here's an exemplary implementation of a NestJS middleware with Express flair:

```
1 import { Injectable, NestMiddleware } from '@nestjs/common';
2 import { Request, Response, NextFunction } from 'express';
3
4 @Injectable()
5 export class DoSomethingMiddleware implements NestMiddleware {
6   use(req: Request, res: Response, next: NextFunction) {
7     /**
8      * Do something here, ...
9      * Chain continues via next() call
10    */
11    next();
12  }
13 }
```

And here's the grand act of wiring middleware into a module (starring the AppModule). Watch closely as the `configure()` method performs its middleware integration magic! `AppModule` must now implement `NestModule` to join this middleware theater - *a mandatory role for any module aspiring to integrate middleware into its performance*:

```
1 import { Module, NestModule, MiddlewareConsumer, RequestMethod } from
  ↳ '@nestjs/common';
2 import { DoSomethingMiddleware } from './...';
3 import { SomeModule } from './...';
4
5 @Module({
6   imports: [SomeModule],
7 })
8 export class AppModule implements NestModule {
9   configure(consumer: MiddlewareConsumer) {
10     consumer
11       .apply(DoSomethingMiddleware)
12       .exclude({ path: 'health', method: RequestMethod.GET })
13       .forRoutes(
14         { path: 'something/*', method: RequestMethod.ALL },
15         { path: 'api/v1/*', method: RequestMethod.ALL },
16         'specific-route'
17       );
18   }
19 }
```

Ignore Import Paths For Now

For our middleware demonstration, we're abstracting away the specific import paths for our custom `DoSomethingMiddleware`. While project structure and naming conventions play their vital roles in the grand production, these paths aren't essential actors in understanding middleware integration. Thus, our imports now gracefully bow out with a simple `from './...'`;